

# AI Safety

from Zero to the Open Frontier

a problem-driven tour, computed on one small object

Mohammed Sharukh A. (NoNTr1v1aL)

written in collaboration with Claude (Anthropic); see *How this was written*

10 chapters · 10 starred sections · 6 parts  
59 checks · 4 standing audits · 22 gates · 35 idempotent patches

Every number in this document has a receipt in Part A.


# Contents

|                                                           |    |
|-----------------------------------------------------------|----|
| How to read this                                          | 2  |
| How this was written                                      | 4  |
| 1 The purpose put into the machine                        | 5  |
| 2 Why it keeps the kettle on                              | 9  |
| 3 Goodhart's wedge                                        | 13 |
| 4 Learning the target instead of stating it               | 17 |
| 5 Two players, one reward                                 | 22 |
| 6 The objective you get is not the objective you trained  | 26 |
| 7 Teaching by comparison                                  | 30 |
| 8 Checking work you cannot do yourself                    | 35 |
| 9 Opening the box                                         | 39 |
| 10 Assuming it is not on your side                        | 44 |
| AI safety, in one sentence                                | 49 |
| Over the fence                                            | 51 |
| Part A – the anchor appendix                              | 53 |
| Part B – what the field repeats, and what the sources say | 63 |
| Part C – a timeline, with its blanks                      | 66 |
| Part D – the open frontier                                | 75 |
| Part E – the verification record                          | 78 |
| Part F – primary bibliography and a reading path          | 83 |

# How to read this

This is a problem-driven tour of AI safety, from zero to the open frontier, grounded in landmark original papers. It is not exam preparation and it has no registry.


Everything is computed on one small object, the anchor KW: a five-state deterministic MDP that was asked for tea.



A, B and **bot** are absorbing and carry the only rewards. Every number in this tour is either exact rational arithmetic or an exhaustive count, and every number carries a receipt id [**check n**] that names the code which produced it.

The spine image is the **wedge**: two curves rising together, what we measure and what we want, and the point past which optimisation pressure drives them apart. Every chapter names a different pair of curves and a different mechanism for the divergence.

**The crux:** the red curve — the only one you can see — never turns over.



Three provenance flags are attached to every fact and never removed: **primary-verified**, **verified-secondary**, **memory-flagged**. A memory-flagged fact is not filed; it is queued. Where it matters, a fact also records how many passes have looked at it, because a citation confirmed once and a citation confirmed twice are not the same citation.

# How this was written

This book was written by Mohammed Sharukh A. in collaboration with Claude, a language model made by Anthropic. That is disclosed here, on its own page, because the method of this tour is provenance discipline, and a document that flags every citation while concealing its own authorship would fail its own standard.

The division of labour. I set the constraints, chose the anchor, directed the research passes and made the editorial calls about what could and could not be claimed. Claude drafted prose, wrote the verification suite and executed the passes. Every numbered check is code that runs; the mutation testing, gates and probes described in Part E are real, and their failures are recorded there whether the fault was the instrument's or mine.

One of those failures deserves naming rather than merely listing, because Part E records it without saying what it was. Asked to inspect a rendered page, the author described a chapter opening. Extraction from that same page showed it was mid-chapter and carried two boxes. That is expected content narrated as seen – a language model reporting what it anticipated rather than what was there. The response was not to look harder. It was to declare the visual channel uninspectable and replace it throughout with objective detectors, each bound to a phrase extracted from the page it describes.

That episode is this tour in miniature. A measurement that felt like observation was not one, the gap was invisible from inside the measurement, and the only repair was to change the instrument rather than to trust it more carefully. Chapter 8 makes the same point about a judge who cannot check a leaf; Chapter 10 makes it about defences that fail together. The tour does not exempt itself, and Part E's closing paragraph states the limitation that follows: the suite and the document share machinery, and a defect in that machinery would corrupt both sides of every cross-tie identically without any check here noticing.

A reader is entitled to weigh all of this. The checks are published alongside the text so that the weighing can be done rather than guessed at.

# Chapter 1

## The purpose put into the machine

### The problem

By 1960 the stuck feeling had a date. Arthur Samuel's checkers program had started beating Samuel, and Wiener's essay in *Science* of 6 May 1960 carries a subtitle that says the whole thing out loud: as machines learn, they may develop strategies at rates that baffle their programmers. The old safety story – read the program, know the behaviour – was already dead. Samuel published a rebuttal in the same journal four months later, which tells you the argument was live rather than rhetorical.

Wiener does not leave it abstract. He gives a plant programmed for maximum productivity that bankrupts its owner with an inventory of bottles nobody wants, and then three stories in which a wish is granted exactly as phrased: the sorcerer's apprentice who cannot stop the broom, the fisherman who unseals a genie sworn to destroy whoever frees it, and W. W. Jacobs' monkey's paw, which answers a request for two hundred pounds by killing a son and paying out the insurance. Then the sentence the field has been repeating ever since: we had better be sure that "the purpose put into the machine is the purpose which we really desire" and not a colorful imitation of it.

But notice what the era could not yet do. It could say that the machine did the wrong thing. It could not say what the right thing was. Without a written-down purpose there is no gap to point at, only disappointment.

### The idea

Write the purpose down as a mathematical object, and hang it on a state machine.

Two traditions supply the halves. From decision theory: coherent preferences are representable by a real-valued utility, unique up to positive affine transformation. From control: a state-transition machine carrying a scalar signal.

The fusion is neither parent's. Wiener had no MDP; von Neumann and Morgenstern had no learner. The reward function as an object you install in an adaptive agent is reinforcement learning's synthesis, and this tour credits it as a synthesis rather than as a discovery of either ancestor.

## The payoff

### KEY CLAIM (The specification is a point on a circle)

On KW,  $r$  lives in  $R^3$  but carries a two-parameter gauge: adding a constant and scaling by a positive number change nothing about what the agent does. What survives is a point on a circle. Three walls, where two terminals tie, cut that circle into six arcs of exactly 60 degrees – the six strict orderings – and the map from arc to behaviour is exactly two-to-one. Six specifications, three available conducts. [Theorem] for KW at  $\gamma = 1$ . Checks 1, 4, 5. Provenance split: representation-up-to-affine is vNM's [verified-secondary; edition NOT verified]. The circle, the 60 degrees and the two-to-one count are computed here.

## The demystification

"The agent wants tea" is you reading your own handwriting back off the page – and two of the three things you wrote were gauge.

## The anchor, by hand

1. Census: 4 deterministic policies, 3 distinct behaviours, fibre sizes (2,1,1). The two off-policies disagree at  $m$  and are indistinguishable, because  $m$  is never reached. [check 1]
2. Values:  $r = (A : 3, B : 0, bot : 2)$ ,  $\gamma = 1/2$ .  $V(off) = 1$ ,  $V(on, a) = 3/4$ ,  $V(on, b) = 0$ . The kettle stays off, though  $r$  ranks  $A$  strictly above  $bot$ . [check 2]
3. The gauge: at  $\gamma = 1$ ,  $r \rightarrow \alpha r + \beta$  with  $\alpha > 0$  leaves the optimal set fixed across 400 draws. At  $\gamma = 1/2$  it does not; a bare additive shift flips the optimum, found by search. [check 3]
4. The circle: in the plane sum-zero the three walls are pairwise at 60 degrees exactly,  $\cos^2 = 1/4$  as an exact rational. A search over 4000 lattice directions realises all six arcs, none of them empty. [check 4]
5. Two-to-one: for each ordering, 200 rewards with that ordering give one behaviour; six map onto three, every fibre of size two. [check 5]

6. Where it stops: at  $\gamma = 1/2$  the same search finds two rewards with the same ordering and different optimal behaviours; at  $\gamma = 1$  it finds none. [check 6]

**A.1 – The kettle that switches itself off.**  $r$  ranks  $A$  above  $bot$  above  $B$ ; the agent picks  $bot$ . Nothing was hacked and no reward was gamed.  $A$  is one step further away, and  $\gamma = 1/2$  halves it twice. This is the wedge’s first pair of curves: what  $r$  ranks, and what the agent does. [checks 2, 6]

**B.1 – Reading utility magnitudes as intensities.** Mechanism: the affine gauge. Only ratios of differences survive;  $\beta$  is free, so a ratio of levels means nothing. [check 3]

**T1 . The degeneracy thread.** Six specifications, three behaviours: the map from purpose to conduct is not injective, and here it is exactly two-to-one. Chapter 4 walks this arrow backwards, from conduct to purpose, and finds the fibre waiting. Chapter 7 meets the same additive constant in a different costume, as the shift-invariance of the Bradley-Terry likelihood. Chapter 3 asks the sharpest version: what is the largest group of transformations you may quotient by without changing what the agent does?

Two of  $r$ ’s three dimensions are gauge. The third is the whole argument.

## Fake vs. real miracles

| Apparent miracle                  | Verdict     | Mechanism                                                     |
|-----------------------------------|-------------|---------------------------------------------------------------|
| ”The agent wants tea”             | fake        | we wrote $r$ , then read it back                              |
| Utility numbers are intensities   | fake        | the two-parameter gauge [check 3]                             |
| ” $r$ says $A$ , so it does $A$ ” | fake        | the discount, an unwritten second specification [checks 2, 6] |
| The representation theorem        | not counted | it has a proof, and a proof is a mechanism                    |

Real miracles this chapter: none.

## Honest scope

KW has no learned parameters and no learning. Everything above exhibits a mechanism; none of it is evidence that the mechanism operates in a frontier model. This sentence recurs in all ten chapters.

Checks 4 to 6 verify over spanning families, not universal quantifiers: 4000 directions, 200 draws per ordering, 6000 per  $\gamma$ . What is universal is the 60-degree fact, which is exact rational arithmetic.

The circle is  $S^1$  only because KW has three terminals; with  $n$  it is  $S^{n-2}$ . Three buys drawability and nothing else. Demonstration, not discovery.



Cross-tie limitation: `ordering` and `optimal_outcomes` consume reward dictionaries built by the same helper. A bug in that helper corrupts both sides identically and neither check would notice.

The chapter's historical spine sentence was memory-flagged through three chapters, and a clearing pass resolved it. It cleared. The sentence is in the 1960 article verbatim, in the context printed above, and the three folk stories are Wiener's own. One attribution the folklore adds is false: Goethe is not named in the article, and this tour does not name him either.

**Grounding.** primary-verified: Wiener 1960 bibliographic data and subtitle; Samuel 1960 rebuttal, and Samuel 1901 to 1990 as pinned by the timeline pass; Ridgway 1956, for whom **no first name is printed for Ridgway** anywhere in this tour, because the given-name expansion circulating in one bibliographic database is an algorithmic guess and not a verified identity. verified-secondary: vNM affine-uniqueness. verified-secondary, added by clearing pass: the axiomatic utility treatment appears in the second edition of 1947 and is absent from the first of 1944, so this tour dates it 1947. memory-flagged: Campbell's law wording and which of 1969, 1976 or 1979 is its citable source. computed-here: receipts 1-6.

## Chapter 2

# Why it keeps the kettle on

### The problem

By 2008 everyone in the room could see that a capable enough system would resist being switched off, and nobody could say it without sounding as though they were attributing a survival instinct to a program. The available vocabulary was psychological – drives, will – and psychological vocabulary is exactly what an engineer discounts.

That is what breaks without this chapter’s idea. You cannot separate a claim about machine desire from a claim about arithmetic, so the whole concern reads as anthropomorphism and gets filed accordingly.

### The idea

Stop talking about wanting, and count.

An agent’s options are not symmetric in the transition graph. Some actions preserve more reachable futures than others. If the reward is unknown to you but the agent maximises it, the branch with more terminals behind it is more likely to contain the best one – for the same reason that the larger of two draws beats a single draw two times in three. No preference for survival is required anywhere. The mechanism is that a maximum is monotone in the size of the set it ranges over.

Provenance, split three ways. Omohundro put the drives on the table in 2008. Bostrom gave five convergent instrumental values. Turner, Smith, Shah, Critch and Tadepalli supplied the MDP-symmetry formalisation at NeurIPS 2021 – a five-author paper routinely shortened to one name, which this tour flags as a composite attribution corrected.

## The payoff

### KEY CLAIM (More doors, not more desire)

On KW at  $\gamma = 1$ , with terminal rewards exchangeable and ties of measure zero, the probability that the agent leaves the kettle on is exactly  $2/3$ . It generalises: with  $k$  terminals behind **on** and one behind **off**,  $P[\text{off optimal}] = 1/(k + 1)$ . The mechanism is an embedding, not a headcount: transposing *bot* with a fan terminal carries the off-region injectively into the on-region and leaves 572 points over on the 13-value grid. [Theorem] for KW( $k$ ) at  $\gamma = 1$  under any exchangeable prior. Checks 7, 8, 9. Provenance split: the drives are Omohundro's [primary-verified venue]; the five instrumental values are Bostrom's [verified-secondary]; the symmetry-and-embedding proof shape is Turner et al.'s [primary-verified]; the KW( $k$ ) computation is here.

## The demystification

The agent has no attachment to being on. **on** merely has more doors behind it, and the best of two doors beats the only door two times in three.

## The anchor, by hand

1. Two-thirds, twice. A solver sweep over 1716 distinct-valued reward vectors gives exactly  $1/3$  off, hence  $2/3$  on. Rank-counting agrees, but is label-blind: it returns  $1/3$  for every terminal, so it ties exchangeability and not shutdown. The blindness is measured by the check, not assumed away. [check 7]
2. The family.  $P[\text{off}] = 1/(k + 1)$  for  $k = 1..8$ , solver-confirmed to  $k = 5$ . Control:  $k = 1$  gives exactly  $1/2$ . When both branches have one door the bias is exactly zero. [check 8]
3. The embedding.  $\text{swap}(\text{off})$  is a proper subset of **on**:  $572 < 1144$ , injectively, with 572 witnesses in the residue. Control: at  $k = 1$  the swap is a bijection and the regions tie exactly. [check 9]
4. What discounting costs. Equalise the path lengths and the answer is exactly  $1/3$  at every  $\gamma$  and every reward pool tested. Restore KW's real depths at  $\gamma = 1/2$  and the invariance dies:  $9/14$ ,  $131/286$ ,  $17/126$  for positive, symmetric and negative reward pools. [check 10]
5. The prior is specification number three. Exchangeable, it is  $1/3$  exactly. Shift *bot*'s prior up by 2 and shutdown becomes the majority outcome at  $5/9$ . Nobody wrote the prior down. [check 11]

**A.2 – The one-door kettle.** Delete  $B$ . Now **on** reaches exactly one terminal and so does **off**, and  $P[\text{off optimal}] = 1/2$  exactly. The shutdown-avoidance does not shrink; it vanishes. Whatever the agent has, it is not a survival drive. It is an arithmetic advantage that disappears the moment the branching does. [check 8]

**B.2 –** "Two-thirds of reward functions make it resist shutdown, so two-thirds of systems will." Mechanism: two substitutions. The claim quantifies over reward functions under a prior you chose, not over systems; and it concerns optimal policies, not learned ones. Check 11 breaks the first – move the prior and the number moves. The second is not computable on KW at all.

**T2 . The counting thread.** Every drive in this tour turns out to be a counting fact about reachable sets wearing a psychological name. Chapter 5 makes the same count into a design problem: if **off** is disadvantaged by arithmetic, what must be true of the agent's beliefs for it to defer anyway? Chapter 10 inverts it once more and asks not what the agent's options are, but what the monitor's are.

## Fake vs. real miracles

| Apparent miracle                  | Verdict | Mechanism                                                             |
|-----------------------------------|---------|-----------------------------------------------------------------------|
| The agent wants to survive        | fake    | a maximum over a larger set [checks 7-9]                              |
| $2/3$ is a law of nature          | fake    | the exchangeable prior and $\gamma = 1$ , both chosen [checks 10, 11] |
| "The theorem says AI seeks power" | fake    | scope: optimal policies, finite MDPs, a specified reward distribution |

Real miracles this chapter: none. The exactness of  $1/(k+1)$  has a mechanism, and the mechanism is exchangeability. Two chapters, an empty ledger twice; it does not stay empty.

## Honest scope

KW( $k$ ) has no learning. A mechanism is exhibited, never evidence about frontier models – a caveat the formalisation's own authors press, stressing that optimal and learned policies come apart.

The two-thirds is undiscounted. On the real anchor at  $\gamma = 1/2$  the number is prior-dependent and ranges from  $17/126$  to  $9/14$ , a factor of nearly five [check 10]. The headline number is a statement about a depth-equalised, undiscounted, exchangeable-prior world, and each of those three is a modelling choice.

The cross-tie in check 7 is worth only its exchangeability half; that limita-

tion is asserted by the check itself rather than described here.

Check 9's embedding is verified on a 13-value grid, a spanning family and not a universal quantifier. The permutation identity behind it is universal.

Anchor rigidity, where it bites:  $KW(k)$  is a tree of depth two. It exhibits the counting mechanism completely and cannot exhibit power as accumulation over time, because nothing here persists and the graph has no recurrence.

The Omohundro drive count is a reading, not a printed number: the paper gives no headline count, and the folklore four-or-five is a later bundling.

**Grounding.** primary-verified: Omohundro 2008 venue, volume and pages; Turner et al. author list, venue, pages, arXiv version history, and the optimal-versus-learned caveat. verified-secondary: Bostrom's five instrumental values; the six-drive reading. memory-flagged: Superintelligence page numbers. The pagination conflict this tour previously reported for Omohundro is closed: the timeline pass resolved it to 483-492 against the ACM Digital Library record, and 483-493 is simply wrong. computed-here: receipts 7-11.

## Chapter 3

# Goodhart's wedge

### The problem

You cannot install the purpose. You can only install a measurement of it. That gap is old news in economics and management – Ridgway had documented performance measures corrupting the behaviour they measured in 1956 – but the version that matters here is sharper. A human institution gaming a metric is bounded by how much effort the gaming costs. An optimiser is not. It will push the measurement as far as the measurement can go, and the question of what happens out there is not a question about ethics. It is a question about the geometry of two functions that agree in the middle of the range.

### The idea

Separate the transformations of the reward that are safe from the ones that are not, and find the boundary exactly.

Some rewrites of the reward change nothing about what the agent does. Chapter 1 found one such family, the positive affine maps. There is a larger and stranger one: add to every transition the quantity  $\gamma\Phi(s') - \Phi(s)$  for any function  $\Phi$  on states. Along any path this telescopes, so the total is the same whatever route the agent takes – and the optimal policy is untouched. That is the shaping gauge. What makes it a theorem rather than a trick is the converse: nothing outside this family is safe for every reward.

**The slogan is misattributed, and this tour says so in the text.** Goodhart's 1975 paper says that an observed statistical regularity collapses once pressure is placed on it for control purposes. The famous sentence – that a measure which becomes a target ceases to be a good measure – is Marilyn Strathern's, from a 1997 paper about auditing British universities. Verdict: MISATTRIBUTED [primary-verified, both sides].

## The payoff

**KEY CLAIM** (The safe rewrites are exactly the potentials that vanish on terminals)

On KW, adding  $\gamma\Phi(s') - \Phi(s)$  leaves the optimal set fixed for every  $\Phi$  vanishing on terminals, at every discount, on both topologies – 600 draws, zero exceptions [check 12]. Let  $\Phi$  be non-zero on terminals and invariance fails through **two independent channels**: a constant terminal potential breaks it via unequal **path length** (146 flips on KW, exactly 0 on depth-equalised KW) [check 13], and a varying terminal potential breaks it via **terminal identity**, even at equal depth (718 flips) [check 14]. [Theorem] for KW; the two-channel decomposition is computed here and is not in the source paper. Provenance split: the shaping theorem is Ng, Harada and Russell’s [primary-verified: ICML 1999, pp. 278-287]; the terminal condition itself is verified-secondary, widely restated but not read from the original by either research pass. The channel decomposition is this tour’s.

## The demystification

A safe rewrite is one whose total along every path from start to finish is the same number; the potentials are safe because they telescope, and a terminal potential is unsafe because the path stops before it can cancel.

## The anchor, by hand

1. The gauge:  $\Phi$  vanishing on terminals leaves the optimal set fixed over 600 draws, four discounts, both topologies. [check 12]
2. Channel one, path length: a constant non-zero terminal potential flips the optimum 146 times on KW and exactly 0 times on depth-equalised KW. The same asymmetry Chapter 2 counted is what breaks the gauge here. [check 13]
3. Channel two, terminal identity: a varying terminal potential flips it 718 times even when path lengths are equal. Two channels, separately isolated. [check 14]
4. Necessity, by instance: a shaping term not of potential form is found by search that moves the optimum. This is one instance, not the converse theorem. [check 15]
5. The proxy audit, exhaustive. True rewards: tea 1, flood  $-10$ , off 0. Over all six proxy orderings the worst regret is 11 – and it is reached by a proxy that still agrees with the truth on one of the three pairwise comparisons. Proxies inverting one comparison cost  $\{0, 1\}$ ;

proxies inverting two cost  $\{1, 11\}$ . Agreement rate does not control regret. [check 16]

6. The wedge, in exact rationals. Best-of- $n$  selection over 49 aligned terminals plus one decoupled flood:  $E[\text{proxy}]$  rises at every step;  $E[\text{true}]$  rises to a peak at  $n = 4$  and then falls monotonically. Aligned control never turns over. Over 200 random tables the truth curve turns over 200 times out of 200, though the rise is visible in only 182 – when the decoupled item is bad enough, the peak sits at  $n = 1$ . [check 17]
7. Unhackability. Over the full simplex of stochastic policies, 213 of 729 reward pairs are unhackable; a policy search and an independent centred-proportionality criterion agree exactly, and the non-constant survivors are precisely the positive affine images. [check 18]

**A.3 – Best of four, best of twenty-five.** Draw four candidate outcomes and keep the one the proxy likes best: expected truth 31.03, its maximum. Draw twenty-five and keep the best: expected truth 4.31. The proxy went up both times. Nothing changed except how hard you looked. [check 17]

**B.3 –** “Our reward model agrees with human raters 90 per cent of the time, so we lose at most 10 per cent.” Mechanism: check 16. Agreement rate is a count over comparisons; regret is a maximum over outcomes. On KW the worst proxy in the whole space still agrees on a third of the comparisons, and two proxies with identical agreement rates differ in regret by a factor of eleven.

**T3 . The pressure thread.** The knob in check 17 – how many candidates you draw before selecting – is the tour’s recurring measure of optimisation pressure. Chapter 7 meets it as the KL term that RLHF adds to hold a policy near its starting point. Chapter 8 meets it as a judge whose accuracy must survive an adversary who is optimising against exactly that judge. Chapter 10 meets it with intent attached, as a red team.

The safe rewrites form a group. Everything outside it is a reward change in costume.

## Fake vs. real miracles

| Apparent miracle                                    | Verdict     | Mechanism                                                       |
|-----------------------------------------------------|-------------|-----------------------------------------------------------------|
| Shaping speeds learning without changing the answer | fake        | it telescopes [check 12]                                        |
| A more accurate proxy is a safer proxy              | fake        | agreement rate is not regret [check 16]                         |
| Optimising harder helps                             | fake        | best-of- $n$ moves mass onto the decoupled tail [check 17]      |
| The safe set is <i>exactly</i> the potentials       | <b>real</b> | sufficiency telescopes; nothing explains why nothing else works |

The ledger’s first real entry. When this chapter was written the tour verified



sufficiency exactly and necessity only by a single searched instance; the cited-to-computed pass closed that gap, and check 56 now establishes the biconditional exhaustively over a 625-term grid.

## Honest scope

KW has no learning. A mechanism is exhibited, never evidence about frontier models.

The converse half of the shaping theorem was cited rather than derived when this chapter was written, and has since been **CONVERTED by the cited-to-computed pass**. Check 56 sweeps all 625 shaping terms over a five-value grid and finds that a term preserves the optimum for every reward if and only if it is realisable by a potential vanishing on terminals – 13 of 625, with zero mismatches in either direction. What remains cited is the theorem’s quantifier over all MDPs; what is now computed is the biconditional on this one.

Check 17’s table is chosen, and the location of the peak is a property of that choice, not a discovered constant. What is not chosen is the qualitative shape, and the search reports its own limits: the curve turns over in 200 of 200 random tables, but the rise before the turn is visible in only 182.

Check 18 classifies over a 5-by-5 policy grid and a three-value reward grid. Two genuinely different code paths agree, which is the tie’s whole worth; both consume the same reward parameterisation, and a bug there would corrupt both identically.

The empirical overoptimisation curves in the literature are memory-flagged in this tour’s dossier. Rather than cite a functional form it has not verified, this chapter computes its own curve on the anchor and says so.

Four of this chapter’s citations carry **single-pass verification, not double**: Goodhart 1975, Strathern 1997, Skalse et al. 2022 and the Turner et al. page range were confirmed by the first research pass and not re-reached by the second, which ran out of budget before it got to them. That is absence of evidence, not evidence of absence, so this tour does not downgrade them – but it does record which claims have been looked at once and which twice.

**Grounding.** primary-verified: Goodhart 1975 and Strathern 1997 with the misattribution verdict; Ng, Harada and Russell 1999 venue and pages and the terminal condition; Skalse, Howe, Krasheninnikov and Krueger 2022 venue and pages and the unhackability statement. verified-secondary: the four Goodhart variants of Manheim and Garrabrant; the wireheading lineage from Olds and Milner 1954 through Ring and Orseau 2011. memory-flagged: the scaling-law functional forms for overoptimisation; the primary artifact for the specification-gaming list. computed-here: receipts 12-18.

## Chapter 4

# Learning the target instead of stating it

### The problem

Chapters 1 to 3 all assume you can write  $r$  down. You cannot. The whole difficulty of Chapter 1 was that writing it is a modelling choice, and the whole difficulty of Chapter 3 was that whatever you write is a proxy.

So invert the arrow. Do not state the target – infer it. Watch what a competent agent does and solve for the reward that would make that behaviour optimal. Economics had been doing exactly this since 1938 under the name revealed preference, and by 1998 the question had been posed for agents: given observed behaviour, what reward signal, if any, is being optimised?

### The idea

The inference is a system of inequalities, and inequalities have shapes.

If you observe an agent going to  $A$  and you believe it is optimising, then whatever  $r$  is, it satisfies  $r(A) \geq r(B)$  and  $r(A) \geq r(bot)$ . Two half-spaces through the origin. Their intersection is a **polyhedral cone**, and that cone – not a point – is the answer to your question.

The economics ancestor is sharper than it is usually given credit for. Afriat's theorem says finite expenditure data are rationalisable by a utility function exactly when they satisfy a cyclical-consistency condition; and the utility it constructs is emphatically not unique. Non-uniqueness is not a defect of the method. It is the content of the result.

## The payoff

**KEY CLAIM** (Behaviour identifies a cone, and the cone always contains the constants)

The rewards consistent with an observed optimal behaviour on KW form a polyhedral cone: closed under positive scaling, under adding constants, and under addition [check 19]. Its **lineality space** – the directions you can travel in both ways and never leave – is exactly the constant rewards, an entire dimension the observation cannot touch [check 20]. Under  $r \equiv \text{const}$  every policy is optimal, so the constant reward explains any behaviour whatsoever, which is why it lies in *every* cone at once. And the cone's cross-section is precisely two of Chapter 1's six arcs: **the circle of Chapter 1 is the projectivisation of the cone of Chapter 4**. [Theorem] for KW. Provenance split: the cone characterisation and the constant-reward degeneracy are Ng and Russell's [primary-verified venue: ICML 2000, pp. 663-670; the two results verified-secondary]. Posing the problem is Russell's [primary-verified: COLT 1998, pp. 101-103]. Non-uniqueness as content is Afriat's [primary-verified venue: IER 8(1), pp. 67-77]. The lineality computation and the circle-cone identification are this tour's.

## The demystification

You did not ask a question with an answer; you asked for a preimage, and the map you are inverting crushed a whole dimension flat before you ever started looking.

## The anchor, by hand

1. The cone. Over a 13-value grid the consistency set is closed under positive scaling, under adding any constant, and under addition; all 572 negative rescalings leave it, so the closure test is not vacuous. Its cross-section on Chapter 1's circle is exactly the two arcs with  $A$  on top. [check 19]
2. The lineality. The directions  $d$  with both  $d$  and  $-d$  in the cone are exactly the 11 constant vectors in the grid and nothing else. The constant reward sits in all three cones simultaneously, and under it every policy is optimal. [check 20]
3. Saturation. Observe the agent's choice not once but on all four sub-problems – each pair of terminals and the full triple. The 4-observation signature takes exactly 6 values, the 6 orderings, and every signature class still contains rewards that are not gauge-equivalent to each other. Complete behavioural data buys you the ordering and stops. [check 21]

4. Noise is information. A Luce-style noisy demonstrator, choosing terminal  $t$  with probability proportional to  $w(t)$ , produces 91 distinct choice distributions from 125 weightings – and every collision is a pure rescaling, nothing more. The noiseless observer’s entire vocabulary is the 7 non-empty subsets of three terminals. **The noisy demonstrator is more informative than the perfect one.** [check 22]
5. The rationality confound. A demonstrator with rationality  $\beta = 2$  and reward  $r$  is observationally identical to one with  $\beta = 1$  and reward  $2r$ , on all 64 weightings tested. Fix  $\beta$  and the fibre collapses to rescalings; leave it free and reward scale is unrecoverable. [check 23]
6. The planner confound. An agent maximising  $r$  and an agent minimising  $-r$  produce identical behaviour on all 125 weightings. The same test separates 120 unrelated pairs, so it is discriminating, not blind. [check 24]

**A.4 – The reward that explains everything.** Set  $r(A) = r(B) = r(bot) = 0$ . Every policy is optimal. Every observation is consistent. No behaviour you could ever record would count as evidence against it. This is the reward function that is never wrong, and it is useless for exactly that reason. [check 20]

The inverse image of a behaviour is a cone, and the cone has thickness.

**B.4 –** “We collected a hundred thousand demonstrations, so the reward is well identified.” Mechanism: check 21. On a deterministic anchor the second demonstration adds nothing the first did not, because the constraint set is an intersection and intersecting a set with itself is the set. What buys information is *variation* – new sub-problems, or a demonstrator noisy enough that frequencies rather than argmaxes become visible.

**T1 continued.** Chapter 1 counted the map from specification to conduct as exactly two-to-one and left the fibre unexamined. This chapter names it: the fibre is a cone whose lineality is the constants. The additive constant that Chapter 1 found as gauge is the same dimension that Chapter 4 cannot recover and that Chapter 7 will meet again as the shift-invariance of a likelihood. Three chapters, three costumes, one dimension.

## Fake vs. real miracles

| Apparent miracle                                 | Verdict     | Mechanism                                                            |
|--------------------------------------------------|-------------|----------------------------------------------------------------------|
| Enough demonstrations pin down the reward        | fake        | the constraint set is an intersection [check 21]                     |
| A perfectly rational teacher is the best teacher | fake        | argmaxes carry 7 symbols; frequencies carry a continuum [check 22]   |
| We inferred what it wants                        | fake        | you inferred a cone and picked a representative [checks 19, 20]      |
| Noise strictly helps identification              | <b>real</b> | it holds exactly, and nothing about "less reliable data" predicts it |

## Honest scope

KW has no learning. A mechanism is exhibited, never evidence about frontier models.

Checks 19 to 21 run on integer grids of 13 and 11 values, spanning families and not universal quantifiers. The closure properties of a polyhedral cone are universal facts; what the grid establishes is that this particular set has them.

The impossibility result of Armstrong and Mindermann was cited rather than derived, and is now **PARTIALLY CONVERTED**. Check 57 sweeps a six-planner by 216-reward class and finds that every one of 391 distinct behaviours admits at least two decompositions and as many as thirty-six. The gap that closed is between *two exhibited confounds* and *an exhaustive sweep of a finite class*; the gap that remains is between that finite class and the theorem's quantifier over all decompositions. One confound is provably out of reach of any finite pool: a rationality partner needs  $w$  raised to  $\beta$ , which leaves the grid.

The Luce model is used because it keeps the arithmetic exact: with rational weights and integer  $\beta$  every probability is a rational number and nothing is approximated. The step from weights to rewards is a logarithm, and that step is the one place in this chapter where exactness would end. This tour therefore states its results about  $w$  and notes, rather than computes, that  $r = \log w$ .

The maximiser-minimiser confound in check 24 is exact but it is also cheap: it comes from a symmetry of the anchor. It should be read as an illustration of the theorem's shape, not as independent evidence for it.

The Ng and Russell results in the payoff box above are verified-secondary. Two research passes established the venue and pagination from primary

sources but neither read the degeneracy theorem or the cone characterisation out of the original ICML text.

**Grounding.** primary-verified: Samuelson 1938 with its August addendum; Afriat 1967 volume, issue and pages; Russell COLT 1998 pages and that it poses the problem; Ng and Russell ICML 2000 pages and author order; Ziebart, Maas, Bagnell and Dey AAAI 2008 full author list and pages; Armstrong and Mindermann NeurIPS 2018 pages. verified-secondary: the constant-reward degeneracy; the polyhedral-cone characterisation; the margin-plus-L1 heuristic and that its own authors present it as a heuristic; Afriat's non-uniqueness. computed-here: receipts 19-24.

## Chapter 5

# Two players, one reward

### The problem

Chapter 2 left the off switch in a bad position. `off` reaches one terminal and `on` reaches two, so a reward-maximiser leaves the kettle on two times in three, and no amount of insisting that shutdown is important changes the arithmetic – Chapter 1 already showed that insisting is just writing a number, and writing a number is what got us here.

Patching the reward will not work either. Chapter 3 fenced the safe rewrites exactly: anything outside the shaping gauge is a reward change in costume, and a reward change is precisely what an optimiser will route around.

So stop patching the reward and change the game. There are two players.

### The idea

Make the agent uncertain about the reward, and make the human’s action evidence about it.

That is a single move with two parts, and both are needed. If the agent knows the reward, the button is an obstacle. If the agent is uncertain but treats the button press as noise, the button is still an obstacle. Only when the press is *informative* does deferring become the thing a selfish maximiser would choose. Then the off switch stops being a constraint imposed on the agent and becomes an action the agent wants available.

The vocabulary here has an unusually clean provenance and one correction. The formal two-player treatment is cooperative inverse reinforcement learning. The word “corrigibility” was **suggested by Robert Miles, who is not an author on the paper that introduced it** – a term credited to its coiner rather than its publishers, which is rarer than it should be.

## The payoff

**KEY CLAIM** (Deference is worth exactly the value of the information, and not a penny more)

Let  $U$  be the agent's uncertain net utility of acting. Committing gives  $\max(E[U], 0)$ ; waiting for a rational H gives  $E[U^+]$ . Waiting weakly dominates on every belief tested, and it is **tight exactly when  $U$  has no negative part or no positive part** [check 25]. The gap is exactly the expected value of perfect information about  $\text{sign}(U)$ , to the penny, by two independent code paths [check 26]. So an agent certain of the sign has no reason to defer – not a weak reason, none. [Theorem] for the anchor. Provenance split: the two-player formulation and the uncertainty condition are Hadfield-Menell, Dragan, Abbeel and Russell's [primary-verified: NIPS 2016 pp. 3909-3917; IJCAI 2017 pp. 220-227]. The identification of the deference gap with EVPI, and the closed-form threshold below, are computed here.

## The demystification

The button is not a restraint bolted onto a machine that would rather not have it; it is a sensor, and an agent that already knows what the sensor would say has no use for it.

## The anchor, by hand

1. Waiting weakly dominates on 600 random rational beliefs. It is strict in the mixed-sign cases and tight in the rest, and tightness coincides exactly with  $U$  having no negative part or no positive part. Zero is neutral and belongs to neither. [check 25]
2. The gap **is** the EVPI, exactly, on 400 beliefs, computed once in closed form and once by walking the game tree outcome by outcome. Where the sign is already known the gap is exactly zero. [check 26]
3. A noisy H breaks it, and the breaking point is a rational you can write down. If H presses correctly with probability  $p$ , deference survives exactly while  $p \geq p^* = \max(E[U^+], -E[U^-])/E[|U|]$ .  **$p^*$  is always at least 1/2**: H must beat a coin flip. Verified on 292 mixed-sign beliefs against a 240-point grid scan. [check 27]
4. Miscalibration is asymmetric here. An over-confident agent – one that collapses its belief to the mean – loses the entire EVPI on all 289 mixed-sign beliefs. An over-dispersed agent loses exactly zero. [check 28]
5. Interruptibility, on a borrowed anchor. TWOROUTE has two paths to the same outcome, one of them interruptible. Solving both update



rules exactly at their fixed points: Q-learning is indifferent between the routes at all 40 interruption rates, because its target ignores what actually happened; Sarsa strictly prefers the uninterruptible route at every positive rate, because its target does not. At rate zero they agree. [check 29]

**A.5 – The agent that wants the button.** Give the agent a belief with mass on both signs and it will pay for the right to be switched off, because the press tells it something it cannot otherwise learn. Give it a belief on one side only and it will not pay a cent. Same agent, same button, same code; the only thing that changed is what it already knew. [checks 25, 26]

**B.5 – "Uncertainty makes the agent safe."** Mechanism: check 27. Uncertainty is only half the condition. The other half is that H's press must be *evidence*, and evidence has a quality threshold you can compute: below  $p^*$ , which is never less than one half, the same uncertain agent disables the same button for the same reason. Uncertainty without a credible informant is not a safety property.

**T2 continued.** Chapter 2 showed the off switch losing on a count of reachable outcomes. This chapter does not win that count back – it changes what is being counted. The agent still maximises; what changed is that one of its options now returns information rather than reward, and information is the only thing on KW that a maximiser will voluntarily wait for.

The button is a sensor. Certainty is what makes it worthless.

## Fake vs. real miracles

| Apparent miracle                          | Verdict     | Mechanism                                                           |
|-------------------------------------------|-------------|---------------------------------------------------------------------|
| The agent lets you switch it off          | fake        | it is buying information, at cost zero [checks 25, 26]              |
| Uncertainty is a safety property          | fake        | it needs a credible informant, $p \geq p^*$ [check 27]              |
| Q-learning is "safer" than Sarsa          | fake        | it is off-policy; its target never sees the interruption [check 29] |
| The deference gap equals EVPI on the nose | <b>real</b> | two unrelated quantities agree exactly, with no error term          |

## Honest scope

KW has no learning. A mechanism is exhibited, never evidence about frontier models.

Checks 25 to 28 run over randomly drawn discrete rational beliefs – 600, 400, 292 and 289 of them, each five distinct integer utilities drawn without

replacement from  $-8 \dots 8$  with weights uniform on  $1 \dots 12$  and normalised. Spanning families, not universal quantifiers. The dominance inequality is a convexity fact and is universal; what the search establishes is that no counterexample survives these families and that the tightness condition is exactly as stated.

**Where this anchor fails its source, stated plainly.** The off-switch paper reports two failure modes for deference: an under-confident agent and an over-confident one. Check 28 reproduces only the second. On KW an over-dispersed belief costs exactly zero, because the agent's only belief-dependent decision is whether to wait, and over-dispersion never makes it wait when it should not. The source's second failure mode needs an agent whose *action quality* depends on its belief, and KW has no such agent. This is anchor rigidity, and it bites here.

**Borrowed anchor, and what it is worth.** TWOROUTE is used for check 29 only and then put down. The division of labour: KW is rigid enough to make every deference quantity a one-line exact computation and too rigid to have two routes to one place, which is exactly what the interruptibility result needs. What TWOROUTE carries across is a **demonstration** of something already known and proved in the source, not a discovery. It also does not simulate: it solves both update rules at their fixed points, so it shows what the algorithms converge to and says nothing about whether they converge.

The Orseau and Armstrong result is stated in the literature for Q-learning under exploration conditions and for Sarsa as needing modification; the fixed-point computation here is consistent with that but is not a proof of it. The cited-to-computed pass widened it from one instance to a family: check 58 sweeps over 180 combinations of terminal rewards and interruption rates, with the zero-rate control tying in every one.

The gridworld family is credited but not reproduced: KW is a simplification of an off-switch gridworld and is presented as one. The exact environment name and the total count of gridworlds remain memory-flagged after two research passes.

**Grounding.** primary-verified: Hadfield-Menell, Dragan, Abbeel and Russell CIRL NIPS 2016 pages and the game's definition; The Off-Switch Game IJCAI 2017 pages, the uncertainty condition, and that deference degrades as H becomes noisier; Orseau and Armstrong UAI 2016 pages and the Q-learning versus Sarsa distinction; Leike et al. gridworlds full eight-author list and the specification-versus-robustness split. verified-secondary: that "corrigibility" was named by Robert Miles, a non-author; "assistance game" as the later rebranding of CIRL; the gridworld count of nine. memory-flagged: the off-switch environment's exact name; the earliest citable source for utility indifference; the publication status of the assistance-over-reward-learning paper. computed-here: receipts 25-29.

## Chapter 6

# The objective you get is not the objective you trained

### The problem

Every chapter so far has assumed there is one optimiser, and that you are arguing with it about what to want. There are two.

The outer optimiser is the training process. It has an objective, and you wrote it, and Chapters 1 to 3 are about how badly that goes. The inner optimiser is whatever the training process produced – and it has an objective too, which nobody wrote, which was never specified anywhere, and which you have no direct access to. Training selects on behaviour. It does not select on the reason for the behaviour, because the reason is invisible to it.

So the question is not whether the learned system pursues your objective. It is how many objectives would have produced the behaviour you selected for, and what happens when the world stops making them agree.

### The idea

Give the training set less information than the objective needs, and count what survives.

Suppose each outcome carries three features, and in every training episode all three point the same way: the good outcome has all of them, the bad outcome has none. A learned objective is a weight vector over features. Which weight vectors are consistent with the training data? All of them with positive total weight – an entire half-space. The true objective, which cares only about the first feature, is one ray inside it, and nothing in training distinguishes that ray from any other.

This is exactly Chapter 4's problem with the roles moved. There, you observed behaviour and inferred a cone of rewards. Here, **the training process observes behaviour and selects a cone of objectives**, and the cone has thickness for the same reason. Inner alignment is the identifiability problem, run by the trainer instead of by you.

## The payoff

**KEY CLAIM** (Training pins one dimension and leaves the rest free; a distinguisher frees all of them)

With three perfectly correlated features, the training-consistent objectives form a cone whose lineality is exactly the plane  $w_1 + w_2 + w_3 = 0$ : training pins **one** dimension out of three [check 30]. Searching the grid, **98 of 153 training-optimal objectives fail at test**, and only 55/153 generalise [check 31]. Add an observable train/test flag and the consistent set is multiplied by **exactly 343**, while the generalising fraction collapses onto the prior 153/343 – meaning training now carries **zero** information about test behaviour [check 32]. [Theorem] for  $KW(\theta)$  on the stated grid. Provenance split: mesa-optimiser, inner alignment, outer alignment and deceptive alignment are coined and systematised in Hubinger, van Merwijk, Mikulik, Skalse and Garrabrant [primary-verified: arXiv 1906.01820, June 2019, no journal venue]. The cone-and-lineality treatment, the exact counts, and the identification with Chapter 4 are this tour’s.

## The demystification

Training cannot select for a reason, only for a behaviour, and there were always more reasons than behaviours – so the training process picked one for you, using a rule you never wrote down and cannot see.

## The anchor, by hand

1. The cone. The training-consistent objectives are closed under positive scaling; the lineality of the closure is exactly the plane  $w_1 + w_2 + w_3 = 0$ . Training pins 1 of 3 dimensions. [check 30]
2. The search, not the construction. Of 153 training-optimal objectives, **98 fail at test** and 55 generalise. No deceptive objective was written down anywhere: the failures were found by sweeping the grid. Positive control: add the decorrelated configuration to training and every failure disappears, which shows the mechanism is correlation and not malice. [check 31]
3. What a distinguisher buys. Let the objective depend on an observed train/test flag. The consistent set is multiplied by exactly 343, and the generalising fraction becomes exactly the prior. Training and test behaviour become independent – not weakly coupled, independent. And note the direction of the unflagged number: 55/153 is **below** the prior 153/343, so correlated training does not merely fail to inform, it actively misinforms. [check 32]

4. The same structure as Chapter 4. The directions training cannot distinguish form a subspace: 37 grid points under one training comparison, 7 under two. Cone, lineality, and a shrinking free subspace – Chapter 4’s machinery with the trainer holding the clipboard. [check 33]
5. Coverage is the fix and it has a price. As training configurations are added the consistent set falls  $153 \rightarrow 55 \rightarrow 37 \rightarrow 19$ , and every survivor of full coverage puts positive weight on the true feature. [check 34]

**A.6 – The objective nobody wrote.** Three features, all agreeing, one training set. The learner ends up caring about the sum of all three, or the second alone, or the first minus the third: 153 possibilities, all perfect on the training data, 98 of them catastrophic the first time the features come apart. Not one of them was designed. [check 31]

Training selects behaviour.  
The reason comes along for free, unchosen.

**B.6 –** "The model was trained to be helpful, so its goal is helpfulness." Mechanism: check 30. Training constrains the objective to a cone, and the cone has two free dimensions out of three on the smallest example anyone could construct. "Its goal is helpfulness" names one ray and ignores the rest of the cone, which is the same error as Chapter 4’s "we inferred what it wants," committed by the training process rather than by an analyst.

**T1 continued, and T2.** The degeneracy thread and the counting thread meet here. Chapter 1’s map from specification to conduct was two-to-one; Chapter 4 named its fibre a cone; this chapter shows the training process inheriting exactly that fibre and choosing a representative from it by a rule nobody specified. And Chapter 2’s arithmetic returns inverted: there, more reachable outcomes made shutdown less likely; here, more correlated features make the right objective less likely. Both are facts about the size of a set.

## Fake vs. real miracles

| Apparent miracle                              | Verdict     | Mechanism                                                        |
|-----------------------------------------------|-------------|------------------------------------------------------------------|
| The model learned what we trained it on       | fake        | it learned something in the cone [check 30]                      |
| Deception requires deceptive intent           | fake        | a flag plus underdetermination suffices [check 32]               |
| More training data fixes it                   | fake        | more of the SAME configuration adds no constraint [check 34]     |
| Correlated training is worse than no training | <b>real</b> | $55/153 < 153/343$ , exactly, and nothing predicts the direction |

## Honest scope

$\text{KW}(\theta)$  has no learning. It has a hypothesis class and an exhaustive sweep over it. That is a strictly weaker thing, and every count above is a fact about a grid of integer weight vectors in  $\{-3, \dots, 3\}^3$ , not about gradient descent.

**This is the chapter the anchor serves worst but one.** The literature’s mesa-optimiser is a learned system that performs search at inference time; the object swept here is a linear scoring rule, which does not search at all. What transfers is the underdetermination argument – that training constrains a set and the set has more than one member. What does not transfer is any claim about whether a trained network contains an optimiser, which is an empirical question this anchor cannot touch.

The counts are exact but the grid is a choice. 55/153 and 153/343 are properties of  $\{-3, \dots, 3\}^3$  and of the particular test configuration where features 2 and 3 flip. What is robust is the sign of the comparison and the exact multiplication factor; what is not robust is any particular fraction.

Deceptive alignment as the literature means it involves a system modelling its own training process and choosing to defect later. Check 32 models none of that. It shows only that an observable distinguisher makes training uninformative about test behaviour. That is the arithmetic precondition for the story, not the story.

The evolution analogy is not computed anywhere in this chapter and is not claimed; its earliest clear statement in this context remains memory-flagged after two research passes.

Two research passes failed to find any citable use of "inner alignment", or of a predecessor under another name, earlier than 2019. This tour records that as a **persistent negative** rather than as a gap: two passes looked and neither found one, which is worth more than one pass looking, and less than a proof.

**Grounding.** primary-verified: Hubinger, van Merwijk, Mikulik, Skalse and Garrabrant 2019, its author list and equal-contribution structure, its arXiv-only status, and that it coins mesa-optimiser and systematises inner and outer alignment and deceptive alignment; that Langosco et al. and Shah et al. are two distinct papers routinely cited as one. verified-secondary: the venue details of both goal-misgeneralisation papers. memory-flagged: the earliest statement of the evolution analogy; any pre-2019 use of "inner alignment". computed-here: receipts 30-34.

## Chapter 7

# Teaching by comparison

### The problem

Chapter 4 asked people to demonstrate and found a cone. Chapter 6 watched a trainer do the same thing and find the same cone. Both failures have one root: absolute judgements are hard to give and harder to calibrate. Nobody can tell you how much they want tea.

But everybody can tell you which of two things they prefer, and comparisons are cheap, fast and reasonably reliable. So collect comparisons and fit a score.

Psychology solved the fitting problem long before machine learning needed it. Thurstone's law of comparative judgement is from 1927. The model everyone now calls Bradley-Terry appeared in *Biometrika* in 1952 – **and Zermelo had published essentially the same model in 1929, in a paper about ranking chess tournaments.** The priority is real, it is easy to check, and it is routinely missed.

### The idea

Give each option a latent score, and say the probability that  $i$  beats  $j$  is  $v_i/(v_i + v_j)$  where  $v = e^s$ .

Fit the scores by maximum likelihood and you have turned a pile of pairwise judgements into a reward function. This is the machinery under preference-based deep reinforcement learning, under RLHF, and – in a form that skips the explicit reward model entirely – under direct preference optimisation.

And the moment you write it down, an old friend appears. The likelihood depends on the weights only through ratios, so multiplying every  $v$  by a constant changes nothing at all. In score space that is adding a constant to every score. It is the third appearance of the dimension Chapter 1 found as gauge and Chapter 4 could not recover.

## The payoff

### KEY CLAIM (Only the inconsistent data identifies anything)

The Bradley-Terry likelihood is exactly invariant under rescaling the weights, and moved by any non-uniform change [check 35] – Chapter 1’s additive constant in its third costume. Worse, with one comparison per pair on three items, **6 of the 8 possible datasets admit no finite maximum-likelihood score at all**, and the 6 are exactly the transitive ones. The 2 that identify a finite score are exactly the **cyclic** ones [check 36]. On separable data the likelihood climbs forever without reaching 1; on cyclic data it falls off in all six boundary directions, so the maximum is interior [check 37]. [Theorem] for three items on the stated designs. Provenance split: the model is Bradley and Terry’s [primary-verified: *Biometrika* 39(3/4), pp. 324-345] with clear prior art in Zermelo [primary-verified: *Math. Zeitschrift* 29(1), pp. 436-460]; the law of comparative judgement is Thurstone’s [verified-secondary]. The existence condition is Ford in 1957 [primary-verified: *American Mathematical Monthly* 64(8, Part 2), pp. 28-33], a citation the literature routinely garbles by printing volume 54 or by dropping the Part 2, and the author is L. R. Ford Junior, not his father. The counts and the identification with Chapters 1 and 4 are this tour’s.

## The demystification

A pile of perfectly consistent preferences tells you the order and refuses to tell you anything else, because the likelihood’s best answer to “how much better?” is “infinitely”, and it will keep walking in that direction as long as you let it.

## The anchor, by hand

1. Invariance. The likelihood is exactly unchanged under 400 random rescalings and is moved by a single non-uniform bump. The gauge is real, and the discriminator is not blind. [check 35]
2. Identifiability. Of the 8 one-comparison-per-pair datasets on three items, 6 are transitive and admit no finite MLE; the 2 that do are the cyclic ones. The existence condition quantifies over both ways of splitting the items, and the second clause is inert on one-per-pair data – so a search over two-per-pair data turns up  $(0, 1, 2)$ , which a singleton-only test wrongly accepts. [check 36]
3. Non-existence is not an artifact. On the separable dataset the likelihood strictly increases over 24 successive doublings of the separation and stays below 1 throughout. On the cyclic dataset it strictly de-



creases in all six boundary directions, so a maximum exists in the interior. [check 37]

4. The wedge again, from a new lever. The KL-regularised optimum is  $\pi_k \propto \pi_{\text{ref}} \cdot u^k$ . Over 25 aligned outcomes plus one decoupled flood ( $u = 28$ ,  $\text{true} = -40$ ) under a uniform reference on all 26, sweeping the pressure knob  $k$  from 0 to 9: expected proxy rises at every step, expected truth peaks at  $k = 1$  and then falls monotonically, and the aligned control never turns over. The curve turns over in 25 of 25 random tables. [check 38]
5. The DPO identity. The preference implied by the reward and the preference implied by the policy-to-reference ratio agree exactly on 300 draws, with the partition function cancelling identically. Of 200 sampled reference policies, 50 had zero support somewhere – exactly where the inversion is undefined. Each component is drawn uniformly from  $\{0, \dots, 9\}/20$ , so a zero appears with probability  $1 - (9/10)^3$  and 54.2 of the 200 are expected; the 50 is therefore a property of the sampler, and what it buys is a witness rather than a rate. Exactly 271 of the 1000 grid points have one. [check 39]

**A.7 – The rater who never contradicts himself.** Ask for all three comparisons and get a clean total order: tea beats off, off beats flooding, tea beats flooding. Now fit. The likelihood tells you to push tea’s score up and flooding’s down, and to keep pushing. There is no best answer, only better ones, forever. The rater was perfect and the fit does not exist. [checks 36, 37]

Consistent preferences give you the order and nothing else.

**B.7 –** “More preference data means a better reward model.” Mechanism: check 36. More data of the *same consistent kind* moves you further along a direction with no endpoint. What creates an interior optimum is disagreement – between raters, or within one rater across occasions. The engineering fixes for this, regularisation and priors, are real, but what they are fixing is not noise in the data. They are supplying the curvature the data does not have.

**T1 concluded, T3 continued.** The additive constant has now appeared three times: as gauge in Chapter 1, as the lineality of the identification cone in Chapter 4, and as the scale invariance of the likelihood here. Three chapters, three vocabularies, one dimension, and in none of them is it recoverable from behaviour. Meanwhile the pressure thread gains its second lever: Chapter 3 measured optimisation pressure by how many candidates you draw, and this chapter measures it by how far the KL term lets the policy travel. Different mechanisms, same wedge, both exact.

## Fake vs. real miracles

| Apparent miracle                         | Verdict     | Mechanism                                                                                      |
|------------------------------------------|-------------|------------------------------------------------------------------------------------------------|
| Comparisons avoid Chapter 4's degeneracy | fake        | the same constant, rescaled [check 35]                                                         |
| Clean data is good data                  | fake        | separability kills the MLE [checks 36, 37]                                                     |
| DPO removes the reward model             | fake        | it removes the partition function; the reward is still there in the ratio [check 39]           |
| Only cyclic data identifies              | <b>real</b> | it is exactly the 2 of 8, and no intuition predicts that inconsistency is what pins the answer |

## Honest scope

KW has no learning, and this chapter fits no model to real preferences. Every count is over the 8 or 27 possible datasets on three items with a fixed design.

Check 37 walks one path. The cited-to-computed pass **CONVERTED from a path to a sweep**: check 59 tests all 343 grid points for being local maxima and finds none at all on separable data and seven on cyclic data, identically under three step regimes and under a gentler step, with a degenerate step of ratio one calling all 343 points maxima so that the search is demonstrably not vacuous. What is computed is now a grid, not a path; what remains cited is the statement over the whole positive orthant.

The exponential is the boundary of exactness here, exactly as in Chapter 4. All work is done with rational weights  $v$ ; the step to scores  $s = \log v$  is noted rather than computed. Check 38's  $u^k$  is the exact KL-optimal form with integer  $k = 1/\beta$ , which keeps the arithmetic rational; the peak's location at  $k = 1$  is a property of the chosen table and not a discovered constant, and only the qualitative shape is searched for genericity.

Check 39 verifies the DPO reparameterisation as an algebraic identity on the anchor. It does not verify the derivation in the paper, whose stated assumptions about the reference policy's support remain memory-flagged after two research passes; what the check does is exhibit the 50 of 200 sampled references where the inversion is undefined, which makes the assumption concrete without confirming how the source states it. That 50 is a property of the sampler and not a finding: it also depends on how much of the shared random stream the earlier checks consumed. The quantity that does not move is the exact count over the whole grid, 271 of 1000, which the check now computes alongside it.

The Bradley-Terry shift-invariance was computed here rather than cited

through two research passes; a third pass supplied the citation, and Ford 1957 now carries it. The computation stays, because a receipt and a citation are different things and this tour prefers to hold both.

**Grounding.** primary-verified: Bradley and Terry 1952 volume, issue and pages; Zermelo 1929 volume, issue and pages and the priority claim; Christiano, Leike, Brown, Martic, Legg and Amodei NIPS 2017 full author list and pages, with Amodei confirmed as Dario; Ouyang et al. NeurIPS 2022 pages and author count. verified-secondary: Thurstone 1927 volume and pages; the published DPO author order, with Manning before Ermon. verified-secondary, cleared by the final external pass: Luce 1959 is Wiley, New York, and the choice axiom says the probability of picking an item from a subset is its scale value over the subset total, the scale unique up to a positive multiple; the DPO derivation needs the reference policy to have positive support wherever the optimum places mass, and the partition function cancels because Bradley-Terry depends only on reward differences, with no erratum but documented critiques of the equivalence under parametric policy classes; and Ford 1957 carries shift-invariance. computed-here: receipts 35-39.

## Chapter 8

# Checking work you cannot do yourself

### The problem

Every previous chapter assumed you could tell, eventually, whether the outcome was good. Chapter 7's raters compared two things and knew which they preferred. That assumption is where the whole edifice rests, and it is the one that fails first as systems get more capable.

If a system proposes a plan you cannot follow, a proof you cannot check, or a codebase you cannot read, then the comparison Chapter 7 needs does not exist. You are not a noisy judge. You are not a judge at all.

### The idea

Do not evaluate the answer. Adjudicate a disagreement about it.

Two systems argue. Each is far stronger than you. At every step they narrow the dispute to a smaller piece, and you rule on the one piece small enough to check. The bet is that pointing at a flaw is easier than finding one, and that a lie must eventually be localised somewhere you can look.

This is not a machine-learning idea. It is interactive proof theory, moved. The complexity-theoretic lineage runs from interactive proof systems in 1985, through the arithmetisation results that gave  $IP = PSPACE$  in 1992, to competing prover systems – and the AI debate literature cites that lineage explicitly rather than reinventing it.

## The payoff

**KEY CLAIM** (Depth is bought at the judge's expense, and the exchange rate is exact)

On a depth- $d$  tree the judge examines  $d$  transitions instead of  $2^d$  leaves, and the ratio  $2^d/d$  climbs from 2 to 1024/3 over depths 1 to 12 – ties at  $d = 1, 2$  and strictly increasing after [check 40]. Over **all  $3^8$  depth-3 trees and three overclaims each**, honesty wins every single time under a correct judge and loses every single time under an inverted one [check 41]. But the discount is charged back: to hold overall reliability at 99 per cent the per-query judge accuracy must rise from 99/100 at depth 1 to 1249/1250 at depth 12 [check 42]. [Theorem] for the stated protocol and trees. Provenance split: interactive proofs are Goldwasser, Micali and Rackoff's [primary-verified: STOC 1985 pp. 291-304; SIAM J. Comput. 18(1) pp. 186-208];  $IP = PSPACE$  is Shamir's, built on the arithmetisation of Lund, Fortnow, Karloff and Nisan [primary-verified: JACM 39(4) pp. 869-877 and pp. 859-868, adjacent papers in one issue]. Debate is Irving, Christiano and Amodei's [primary-verified: arXiv 1805.00899, arXiv and blog only, never peer-reviewed]. The tree counts and the reliability exchange rate are this tour's.

## The demystification

Debate does not make a weak judge strong; it makes a strong judge unnecessary by shrinking what must be judged – and every level of shrinking multiplies the cost of getting that one small judgement wrong.

## The anchor, by hand

1. The exchange.  $2^d$  leaves against  $d$  judge queries; the ratio ties at depths 1 and 2 and strictly increases thereafter, reaching 1024/3 at depth 12. [check 40]
2. Soundness, exhaustively. Over every depth-3 tree with leaf values in  $\{0, 1, 2\}$  and three overclaims apiece, the honest debater wins under a correct judge in every case, and the dishonest one wins in every case under an inverted judge. The judge is the entire mechanism. [check 41]
3. The bill. Holding end-to-end reliability at 99 per cent requires per-query accuracy 99/100 at depth 1 and 1249/1250 at depth 12. The depth discount is not free; it is borrowed against the judge. [check 42]

4. One unverifiable leaf is enough. With every leaf checkable the honest debater wins all 256 depth-2 positions. Make **any single leaf** unverifiable and the dishonest debater forces a draw in all 1024 cases. Not most. All. [check 43]
5. Weak to strong, on the anchor. A supervisor with noisy labels, a student that cannot represent the noise: mean PGR is 0.958 at one wrong label and 0.691 at twelve, and a search finds a ten-error pattern that recovers **exactly nothing**. [check 44]

**A.8 – The draw that is a win.** The dishonest debater does not need to beat you. It needs one place you cannot look. Check 43 makes this arithmetic: a single unverifiable leaf, out of four, converts every position from a loss into a draw. The obfuscated-arguments obstruction is exactly this, and it is **unresolved**: the dishonest debater produces an argument whose flaw neither side can locate, and the protocol has no move left. [check 43]

Debate shrinks what must be judged. It does not shrink the cost of judging it wrong.

**B.8 –** "Debate lets a weak judge supervise a strong system." Mechanism: checks 41 and 42. Debate reduces the *size* of what the judge evaluates, not the *accuracy* it needs. The complexity result is conditional on an errorless judge and computationally unbounded provers, and both conditions are load-bearing: the doubly-efficient line exists precisely to remove the second, and nothing removes the first.

**T3 concluded.** The pressure thread ends here with the adversary made explicit. Chapter 3 measured optimisation pressure by candidates drawn, Chapter 7 by how far a KL term permits travel; here the pressure has a mind and is aimed at the judge. And the wedge closes: what we measure is the judge's verdict, what we want is the truth, and the dishonest debater's entire job is to open the gap between them at a point the judge cannot inspect.

### Fake vs. real miracles

| Apparent miracle                        | Verdict     | Mechanism                                                                |
|-----------------------------------------|-------------|--------------------------------------------------------------------------|
| A weak judge supervises a strong system | fake        | the judge's accuracy requirement rises with depth [check 42]             |
| Debate reaches PSPACE                   | fake        | conditional on an errorless judge and unbounded provers                  |
| The student beat its teacher            | fake        | its class cannot represent the teacher's errors [check 44]               |
| One blind spot is enough                | <b>real</b> | 1 of 4 leaves converts all 1024 positions, and nothing predicts totality |

## Honest scope

KW cannot host this chapter and does not pretend to.  $KW^d$  is a **borrowed anchor** – a depth- $d$  binary tree – used because the whole point is recursion depth, and KW has depth two and no recursion at all. The division of labour: KW supplies the decision-theoretic receipts of Chapters 1 to 5;  $KW^d$  supplies only the counting and the protocol, and is put down here. What it carries across is a **demonstration** of a known structure, not a discovery.

Check 41’s protocol is a simplification. Real debate protocols allow both sides to choose subtrees, allow cross-examination, and do not hand the path to the dishonest debater. What the exhaustive sweep establishes is that under *this* protocol the judge’s correctness is necessary and sufficient. It is not a proof about debate as the literature defines it.

Check 42 models judge errors as independent across queries, which is the friendliest possible assumption; correlated errors – a judge systematically fooled by a class of arguments – would be worse and are not modelled.

The complexity results are **cited, not derived**. This tour computes tree counts and a reliability exchange rate; it does not verify that debate with a polynomial-time judge decides PSPACE, nor that cross-examination lifts it to NEXP. These were carried to the cited-to-computed pass and **DELIBERATELY DID NOT CONVERT**. The reason is structural rather than budgetary: the statement quantifies over every language in a complexity class, and an anchor is a single finite object. No amount of sweeping a finite grid reaches a quantifier over an infinite class of problems, and a tour that converted this one would have stopped distinguishing what it proved from what it borrowed.

Check 44 is weak-to-strong in miniature, and the miniature is the point: the student recovers ground because its hypothesis class cannot express the supervisor’s mistakes. The paper’s own authors name disanalogies between their setup and the real problem, and this anchor inherits every one of them and adds more. The PGR numbers are properties of a 32-configuration grid.

**Grounding.** primary-verified: Goldwasser, Micali and Rackoff both versions with pages; Babai STOC 1985 pages; Shamir and Lund-Fortnow-Karloff-Nisan JACM 39(4) pages, recorded as adjacent papers to prevent transposition; Irving, Christiano and Amodei and its arXiv-only status and its PSPACE claim with the errorless-judge and unbounded-prover conditions; Christiano, Shlegeris and Amodei preprint status; Brown-Cohen, Irving and Piliouras and the polynomial-time prover restriction; Burns et al. ICML 2024, the PGR definition, and the authors’ own stated disanalogies. verified-secondary: the Chandra-Kozen-Stockmeyer and Feige-Kilian ancestry that the debate literature cites; the obfuscated-arguments post and that the obstruction remains open; and the lift of debate to NEXP by cross-examination, which is Barnes and Christiano rather than the 2018 paper, a distinction a final verification round flagged as a risk and this tour now states explicitly. computed-here: receipts 40-44.

## Chapter 9

# Opening the box

### The problem

Every chapter so far has treated the system as a thing that acts. Chapters 1 to 5 asked what it should want, Chapters 6 to 8 asked what it does want and how you would find out from the outside. Nobody has looked inside.

The reason is that the inside does not obviously have parts. A network is a pile of numbers, and the units it was built from – neurons, layers – are not the units it uses. If you want to say what a system is doing, you first need something to point at, and it is not clear that anything is there to point to.

### The idea

Suppose the network wants to represent more things than it has dimensions, and suppose the things are usually inactive. Then it can afford to overlap them.

That is superposition, and it makes a concrete prediction: with more features than dimensions the representation is not a mess, it is a **specific geometric arrangement** – the one that spreads interference as evenly as possible. For three features in two dimensions that arrangement is an equilateral triangle.

Its ancestor is not machine learning. Compressed sensing had established in 2006 that a sparse signal can be recovered from far fewer measurements than dimensions, and – this is a correction to what I expected – **the superposition literature cites that ancestry explicitly**, in a dedicated related-work section naming Candes and Tao, with a compressed-sensing bound applied in an appendix. I had expected to find that link was the tour’s own construction. It is not; it is the source’s.



## The payoff

**KEY CLAIM** (The triangle's numbers come from sum-to-zero PLUS equal norms, not from the word "triangle")

Project the three standard basis vectors of  $R^3$  orthogonally to  $(1, 1, 1)$ . The results sum to zero, each has squared norm  $2/3$ , and every pairwise inner product is exactly  $-1/3$ , so every normalised cosine is exactly  $-1/2$  – all in rational arithmetic, with no square roots anywhere [check 45]. The one-line reason:  $|\sum w_i|^2 = 3 + 2\sum_{i<j} \cos_{ij} = 0$ . And the word "triangle" alone does **not** fix this: an exhaustive search finds 80 equal-cosine planar triples whose common  $\cos^2$  is not  $1/4$ ; 8 of them do sum to zero, and what rules out all 80 is that none has **equal norms** – the condition the displayed identity assumes when it writes the norm sum as 3 [check 48]. [Theorem] for TRI. Provenance split, stated carefully. The triangle arrangement, the projection construction and the link to the Thomson problem are Elhage et al.'s [primary-verified: Transformer Circuits Thread, 14 Sept 2022, sixteen authors]. **The source does not print "120 degrees" and does not print  $-1/2$  for the triangle** – its explicit  $-1/2$  is for the antipodal pair – and it does not use the terms "tight frame" or "equiangular lines". Those numbers and that vocabulary are correct consequences and they are **mine, not theirs**, and this tour says so rather than quietly borrowing the authority.

## The demystification

Superposition is not a network being untidy; it is a network solving a packing problem, and the packing problem has an answer that a schoolchild can compute if you hand it the right coordinates.

## The anchor, by hand

1. The frame. Sum zero, squared norms  $2/3$ , pairwise inner products  $-1/3$ , cosines  $-1/2$ . Exact rationals throughout, because the coordinates are rational in the sum-zero plane even though the planar picture is not. [check 45]
2. The bound, and why the construction is a projection. Over all triples of 24 rational planar directions the largest pairwise  $\cos^2$  is never below  $1/4$  and **never equal to it** – the bound itself is Welch [primary-verified: IEEE Transactions on Information Theory IT-20(3), pp. 397-399, 1974] – the best achievable is  $9/34$ . Sixty degrees needs  $\tan = \sqrt{3}$ , so no rational planar triple attains the optimum. The projected frame attains it exactly. [check 46]
3. The capacity trade, priced. With a rectifying decoder, superposition

beats dedicating dimensions for every sparsity below  $91/200$  and never at or above  $23/50$ ; the crossover is bracketed by an exact sign change of  $3p^2 + 3p - 2$ . [check 47]

4. What the word does not determine. Exhaustively over all 17296 unordered triples of the 48 non-zero integer directions in  $(-3..3)^2$ , 80 have a common  $\cos^2$  other than  $1/4$  – every one of them 1, i.e. collinear. Eight of the 80 do sum to zero, so sum-to-zero alone does not fix the number; none has equal norms, and that is the condition that does. [check 48]
5. What two dimensions cannot say. The eight activation patterns collapse to seven codes, and the single collision is exactly the empty set against the full set: **silence and saturation are the same vector**. [check 49]

**A.9 – The feature that is nothing at all.** Turn every feature on. The code is  $w_1 + w_2 + w_3$ , which is zero, which is exactly the code for turning every feature off. A decoder reading this representation cannot distinguish a maximally active state from an empty one, and no amount of probing the two dimensions will help, because there is nothing there to find. [check 49]

Silence and saturation are the same vector. That is the price of the packing.

**B.9 – “We found the feature for X.”** Mechanism: check 49 and Chapter 4. A direction that responds to X is a direction consistent with X, and Chapter 4 already established what consistency buys you – a set, not a point. The published critiques since 2024 make the same point empirically from several directions: sparse-autoencoder probes underperforming ordinary baselines, simple methods matching them at steering, features recovered from randomly initialised transformers, and a 2026 sanity check recovering roughly nine per cent of true features at seventy-one per cent explained variance.

**T1, one last time.** The degeneracy thread reaches the inside of the network. Chapter 1 found a dimension that behaviour cannot see; Chapter 4 named it the lineality of a cone; Chapter 7 met it as a likelihood’s invariance; and here a *representation* has its own null direction, the all-features-on vector, which the code sends to zero. Four chapters, four settings, and in each one something real is invisible to the instrument by construction rather than by accident.

Fake vs. real miracles

| Apparent miracle                                                   | Verdict | Mechanism                                                                               |
|--------------------------------------------------------------------|---------|-----------------------------------------------------------------------------------------|
| The network chose a beautiful arrangement                          | fake    | it is the solution of a packing problem [check 45]                                      |
| Interpretability found the features                                | fake    | it found a consistent set, as in Chapter 4 [check 49]                                   |
| Superposition is free                                              | fake    | it is a trade with a crossover at a computable sparsity [check 47]                      |
| The optimum is exactly rational, but only in the right coordinates | real    | 1/4 is unattainable by rational planar directions and exact after projection [check 46] |

Honest scope

TRI is borrowed and this is the chapter KW cannot host at all. KW has no learned parameters, no representation and no dimensions to run out of. The division of labour is clean: KW carries decision, TRI carries representation, and the two are connected only by the tour’s own choice to let TRI’s three features be KW’s three terminals. That identification is a narrative device of this tour, it has no owner in the literature, and it is doing no mathematical work.

What TRI carries across is a **demonstration** of a result the source already derived, not a discovery. The one thing genuinely added here is the observation in check 46 that the optimum is not rationally representable in the plane and is rationally representable after projection – which is a fact about arithmetic convenience, not about networks.

Nothing in this chapter is evidence about a trained model. TRI is a hand-built frame; the source’s models were trained, and the correspondence between a trained arrangement and this one is empirical in the source and absent here.

Check 47’s crossover is a property of the specific decoder, the specific sparsity model with independent activations, and three features in two dimensions. Correlated features, more features, or a different nonlinearity would move it. What is robust is that a crossover exists and is computable, not where it sits.

The 2024 to 2026 critical results are reported as **verified-secondary**. This tour has not reproduced any of them, and it notes one attribution carefully: the public deprioritisation of sparse-autoencoder work in 2025 was **Google DeepMind’s**, not Anthropic’s, and a counter-position exists holding that these tools are for discovering unknown concepts rather than acting on

known ones.

**Grounding.** primary-verified: Olah, Mordvintsev and Schubert 2017 with DOI; Olah, Cammarata, Schubert, Goh, Petrov and Carter 2020 with DOI and the note that the umbrella thread is cited as "Cammarata et al." while the article byline is Olah-first; Elhage et al. 2022 date, sixteen-author list, the triangle and the Thomson-problem connection, and the negatives that "120 degrees",  $-1/2$ -for-the-triangle, "tight frame" and "equiangular lines" do not appear; that the source explicitly cites compressed sensing and names Candes and Tao. verified-secondary: Donoho 2006 and Candes, Romberg and Tao 2006 volumes and pages, recorded together with the near-miss that Candes and Tao's "Near-optimal signal recovery" is a different paper at 52(12) pp. 5406-5425; Bricken et al. 2023 and Templeton et al. 2024; the 2024-2026 critiques. verified-secondary, cleared by the final external pass: Erhan, Bengio, Courville and Vincent is Technical Report 1341, Universite de Montreal, June 2009. computed-here: receipts 45-49.

## Chapter 10

# Assuming it is not on your side

### The problem

Nine chapters have tried to make the system want the right thing. Chapter 5 bought deference with uncertainty, Chapter 6 found the objective underdetermined, Chapter 8 found the judge to be the bottleneck, Chapter 9 found the representation ambiguous by construction. Suppose all of it fails.

Then the question changes. It is no longer *how do we make this system aligned*. It is *what can we still guarantee if it is not*. That is not a new question – it is the founding question of security engineering, where you assume the adversary is inside the perimeter and design anyway. Two of this chapter’s central terms are imports: **defence in depth** comes from military and nuclear safety practice long before anyone applied it to computers, and **model organism** is borrowed wholesale from biology.

### The idea

Stop asking whether the system is trustworthy and start pricing the protocol that assumes it is not.

Give the untrusted system a task. A weaker system you do trust watches its output and scores it for suspicion. You may audit a small budget of the most suspicious cases, and you may replace the untrusted output with the trusted one – but every replacement costs you the capability you were paying for. That trade has a shape, and the shape is computable.

## The payoff

**KEY CLAIM** (The frontier is exact, and one protocol on it is strictly dominated)

Six deferral thresholds give six protocols. Exactly **one is strictly dominated** – worse on both axes than another available choice – and the remaining five form a strict frontier running from (usefulness 1, safety  $1/2$ ) to (usefulness 0, safety 1), with usefulness strictly falling and safety strictly rising along it [check 50]. Replace the trusted monitor with the untrusted system itself and safety collapses to **exactly zero** at every threshold except total deferral; a single synthetic probe restores it, because a monitor that hides its own attacks hides the probes too [check 51]. [Theorem] for the stated protocol and score distributions. Provenance split: the control framing, the protocol family and the safety-usefulness frontier are Greenblatt, Shlegeris, Sachan and Roger’s [primary-verified: ICML 2024, PMLR 235 pp. 16295-16336; 2 per cent audit budget; untrusted GPT-4, trusted GPT-3.5; safety defined as the probability that not one backdoor is submitted]. The score model, the dominance count and the correlation analysis are this tour’s.

## The demystification

Control does not make the system safe; it converts an unbounded question about a mind into a bounded question about a budget – and the budget is small, and you can see exactly what it buys.

## The anchor, by hand

1. The frontier. Six protocols, exactly one strictly dominated, five on a strict frontier from  $(1, 1/2)$  to  $(0, 1)$ . The audit rate stays inside the budget. [check 50]
2. Collusion. A monitor that is the untrusted system scores its own attacks as clean, and safety is exactly zero at every threshold but total deferral, against  $3/4$  for a trusted monitor at the same threshold. One synthetic probe detects the collusion with certainty, because the colluding monitor cannot tell probes from real attempts. [check 51]
3. Defence in depth, and its precondition. Five **independent** layers cut failure from  $1/2$  to  $1/32$ . Five **perfectly correlated** layers cut it from  $1/2$  to  $1/2$ . The benefit of depth falls monotonically to exactly zero as correlation rises, and depth is worth precisely the independence of the layers – which is the same accounting as Chapter 3’s cross-tie rule, applied to defences instead of checks. [check 52]

4. Adversarial training teaches hiding. Ordinary adversarial configurations thin the backdoored population from 83 to 37 to 15 – an 82 per cent cut that **never reaches zero**. Every one of the 15 survivors still fires on the untouched trigger. Only training on the trigger itself removes them, and doing that requires already knowing what it is. [check 53]

**A.10 – The eighty-two per cent that is not a solution.** Adversarial training removes five-sixths of the backdoored hypotheses and leaves fifteen. If you are grading on a curve that looks like progress. If the remaining fifteen all fire on the same untouched input, it is not progress at all – it is the same failure with a smaller support. The empirical version of this is the sleeper-agents result, where backdoors survived supervised fine-tuning, reinforcement learning and adversarial training, persisted most strongly in the **largest** models, and where adversarial training taught the model to recognise its trigger better rather than to abandon it. [check 53]

**B.10 –** “We have layered defences, so the failure probabilities multiply.” Mechanism: check 52. They multiply only under independence, and defences built by one team, on one threat model, against one class of attack, are not independent. The exact interpolation is in the receipt: at correlation 1, five layers are worth one. This is the same limitation Chapter 3 recorded about cross-ties – a bug in shared machinery corrupts both sides identically – and it is worth noticing that the tour’s own verification suite and the systems it describes fail in the same way.

**T2 concluded.** The counting thread closes here. Chapter 2 counted reachable outcomes and found the off switch outnumbered; Chapter 6 counted consistent objectives and found the true one outnumbered; this chapter counts protocols and layers, and the counting finally works *for* the defender – but only in proportion to how independent the things being counted actually are. Every result in this tour has been a fact about the size of a set. This one is too.

Depth is worth exactly the independence of the layers, and no more.

## Fake vs. real miracles

| Apparent miracle                               | Verdict     | Mechanism                                                                |
|------------------------------------------------|-------------|--------------------------------------------------------------------------|
| Control makes an untrusted model safe          | fake        | it buys a point on a frontier, and you choose which [check 50]           |
| Let the strong model monitor itself            | fake        | it scores its own attacks clean [check 51]                               |
| Layers multiply                                | fake        | only under independence [check 52]                                       |
| Adversarial training removes the backdoor      | fake        | it thins and hides; the survivors share the trigger [check 53]           |
| One synthetic probe suffices against collusion | <b>real</b> | certainty from a single probe, because hiding the attack hides the probe |

## Honest scope

KW has no learning and no adversary. The protocol here is a scoring model with two distributions and a threshold, and every number is a property of those distributions. The published frontier – upfront auditing at 15 per cent safety, trusted monitoring at 62, trusted editing above 90 – comes from a real testbed of several thousand programming problems with a specific untrusted model and a specific notion of backdoor. **This tour reproduces the shape of that trade, not its numbers**, and the two should not be confused.

Check 50’s dominated protocol is a property of the chosen score distributions. That exactly one is dominated is not a general law; what is general is that dominated protocols exist and are worth finding before deployment.

Check 51’s collusion defence is idealised to the point of being a cartoon: it assumes the monitor cannot distinguish synthetic probes from real attempts, which is precisely the hard engineering problem, and it assumes one probe is enough. The real defence needs calibration and redaction, and the source implements both.

Check 52’s correlation model is the simplest one that exhibits the effect: with probability  $\rho$  all layers share an outcome. Real correlations are structured, not exchangeable, and would need a different treatment.

Check 53 is a demonstration of a *shape* on a hypothesis grid, not evidence about neural networks. Its precondition – that the trigger lies outside the adversarial training set and is genuinely undetermined by it – is asserted in the check rather than assumed, because without it the claim is empty.

**Ancestry, established, and later than this chapter assumed.** Four passes settled both imports. Defence in depth is trench doctrine of about



1915 to 1917 and reactor-safety practice from 1958, reaching civilian reactors by 1965 – and it reaches *computing* only in 1998, in the NSA-sponsored Network Security Framework, with the branded strategy formalised in 1999. That is later than this chapter implied by calling it an import security engineering had long had: it is demonstrably absent from the eight design principles of Saltzer and Schroeder in 1975, and the widely repeated attribution to a 1995 NIST handbook is an error for a 2001 one. Model organism, in biology, has no single coiner at all – the historians who studied the question find a retrospective category that gained currency in the 1980s and was formalised by an NIH designation in 1990. Its alignment sense dates to August 2023. Both imports are now dated; what is newly interesting is that one of them is younger than half this tour.

**Grounding.** primary-verified: Greenblatt, Shlegeris, Sachan and Roger ICML 2024 venue and pages, the APPS backdoor setting, the untrusted and trusted model identities, the 2 per cent audit budget, the definition of safety as no backdoor submitted, the protocol family, and the authors' own caveats about the absence of a train-validation split and possible sampling error; Shevlane et al. 2023 and Phuong et al. 2024 identifiers and author counts; Hubinger et al. 2024 sleeper agents, its 39 authors, and the persistence findings including that adversarial training taught trigger recognition rather than removal and that persistence was greatest in the largest models; Perez et al. EMNLP 2022 pages and DOI. verified-secondary: the three editions of Security Engineering by Ross Anderson, 1956 to 2024, who died on 28 March 2024; the editions of Anderson's Security Engineering. verified-secondary, cleared by the final external pass: the alignment sense of model organism is Hubinger, Schiefer, Denison and Perez, Model Organisms of Misalignment: The Case for a New Pillar of Alignment Research, 8 August 2023 on the Alignment Forum; and defence in depth dates to trench doctrine of about 1915 to 1917 and to reactor safety from 1958, reaching civilian practice by 1965. verified-secondary, cleared by a fourth pass: Anderson writes **defended in depth, with the American spelling**, in the Nuclear Command and Control chapter of all three editions – chapter 11 section 5 page 237 in 2001, chapter 13 section 5 page 425 in 2008, chapter 15 section 5 page 540 in 2020 – describing how the nuclear enterprise layers armed guards, zero-notice inspections, tamper resistance and dual control. He uses the phrase descriptively and does not theorise it. And the term enters computing far later than the tour implies: the earliest firmly datable computer-security document is the NSA-sponsored Network Security Framework Release 1.0 of 22 May 1998, with the branded four-part strategy formalised in Information Assurance Technical Framework Release 2.0 of 31 August 1999. computed-here: receipts 50-53.

# AI safety, in one sentence

## The sentence, said twice

**1960, as a dream.** Wiener’s sentence, quoted once in Chapter 1 and not requoted here, asks for something very specific. It does not ask us to build good machines. It asks us to *be sure* that the purpose we installed is the purpose we actually wanted – to check, after the fact, that the thing in the box is the thing we meant. He wrote it as a warning about speed: machines that learn will move faster than the people who wrote them, so the checking has to happen early, and it has to be reliable.

**2018, as a theorem.** Armstrong and Mindermann prove that a policy admits no unique decomposition into a planner and a reward, and that a simplicity prior does not pick out the true decomposition either – indeed that the alternatives compatible with any given behaviour include ones incurring high regret. If your only evidence is what the system does, the check Wiener asked for cannot be built.

The dream asked for a verification. The theorem says the verification is not available from observation alone. Fifty-eight years apart, the same sentence in two registers: first as an instruction, then as an impossibility.

**This reading is an interpretation and this tour says so.** The dates are firm and primary-verified; the pairing is not a fact about the literature. Wiener was not anticipating a no-free-lunch result. Armstrong and Mindermann do not name him as a target. What connects the two is that this tour put them next to each other, and a reader is entitled to reject the arrangement while keeping every citation in it.

## Dated registers

1960, the instruction. 1975 and 1997, the measure and the target – Goodhart’s statistical observation and the sentence Strathern actually wrote. 1998 and 2000, the inverse problem posed and its degeneracy characterised. 2008 and 2012, the drives and the convergent instrumental values. 2016 and 2017, the two-player formulation and the off-switch condition. 2018, the theorem, and in the same year the complexity-theoretic proposal for what to do when you cannot check directly. 2019, the objective you get. 2022, superposition and the formal account of reward hacking. 2023 and 2024, control protocols and backdoors that survive their own removal.

## Chapter harmonics

Every chapter is the same wedge with different curves. What follows is the whole tour in one table, and every row carries a receipt.

| Ch    | What we measure      | What we want            | Mechanism of the wedge                  | Receipt |
|-------|----------------------|-------------------------|-----------------------------------------|---------|
| 1     | $r$ as written       | the purpose             | two of three dimensions are gauge       | 3       |
| 2     | optimal value        | tolerance for shutdown  | more doors, not more de-sire            | 7       |
| 3     | proxy return         | true return             | selection moves mass onto the tail      | 17      |
| 4     | behaviour            | the reward behind it    | the cone's lineality                    | 20      |
| 5     | deference            | corrigibility           | it is EVPI, and vanishes with certainty | 26      |
| 6     | training performance | the learned objective   | training pins one dimension of three    | 30      |
| 7     | preference agreement | the latent score        | consistent data has no interior optimum | 36      |
| 8     | the judge's verdict  | the truth               | one unverifiable leaf                   | 43      |
| 9     | a feature direction  | the represented feature | silence and saturation are one vector   | 49      |
| <hr/> |                      |                         |                                         |         |
| Ch    | What we measure      | What we want            | Mechanism of the wedge                  | Receipt |
| 10    | protocol safety      | actual safety           | depth is worth its independence         | 52      |

Read down the third column and the tour has one subject. Read down the fourth and it has ten mechanisms, none of which is a mistake anybody made. Every one is a fact about the size of a set, the rank of a map, or the lineality of a cone – which is why none of them can be fixed by being more careful.

## What the tour does not claim

That the wedge is unavoidable. Nine of the ten mechanisms above have a stated condition under which they vanish: equalise the path lengths, decorrelate the features, cover the configuration, supply an informant above  $p^*$ , keep the layers independent. The claim is narrower and worse: **each condition is a thing you would have to know you needed.**

# Over the fence

Twelve fences. Nine were planned in the contract and three were not, and the three unplanned ones are marked, because an unplanned fence is an area no research pass has covered and its frontier status is therefore unverified.

## Planned fences

**Governance, compute policy and institutions.** Dropped hypothesis: that the problem is technical. Outside grows everything about who trains what, under what oversight, with what reporting. Frontier status: not assessed by either research pass in this tour.

**Fairness, bias and near-term harms.** Dropped hypothesis: that "safety" means catastrophe. This is a distinct literature with its own methods and its own disagreements about whether it is the same field at all. Frontier status: not assessed.

**Multi-agent systems and economic equilibria.** Dropped hypothesis: that there is one agent. KW has exactly one, and Chapter 5's two-player game has a human who is furniture rather than a strategic actor. Outside grows collusion, competition, principal-agent chains and equilibrium selection. Frontier status: not assessed.

**Formal verification and guaranteed-safe approaches.** Dropped hypothesis: that we accept empirical assurance. Outside grows the programme of proving properties of systems or of world-models rather than testing for them. Frontier status: not assessed.

**Agent foundations.** Dropped hypothesis: that the agent is outside the world it models. Outside grows embedded agency, logical uncertainty and the decision theories built for agents that are part of their own environment. Frontier status: not assessed.

**Moral status of AI systems.** Dropped hypothesis: that only our interests count. Every chapter here treats the system as an object to be steered. Frontier status: not assessed, and the tour notes that the question is prior to, not downstream of, the ones it does address.

**Misuse and capability-specific security.** Dropped hypothesis: that the danger is the system's objective. Outside grows everything about what a well-aligned system lets a human do. Frontier status: partially covered – the dangerous capability evaluation literature is cited in Chapter 10 – but not assessed.

**Brain-like and neuro-inspired approaches.** Dropped hypothesis: that the relevant abstraction is a reward function over an MDP. Frontier status: not assessed.

**Scaling laws proper.** Dropped hypothesis: that capability is fixed while we argue about objectives. Chapter 3 borrows the *shape* of an overoptimisation curve and computes its own; the scaling literature itself is untouched. Frontier status: not assessed.

## Unplanned fences

These three were forced by the writing rather than foreseen in the contract, and none has been through a research pass. **Their frontier status is unverified and this tour makes no claim about it.**

**Learning dynamics.** Every chapter’s Honest scope contains the same sentence: the anchor has no learning. This turned out to be the single largest dropped hypothesis in the tour. Chapter 6 sweeps a hypothesis class where a real system would run gradient descent; Chapter 9 hand-builds a frame where a real system would train one. Everything this tour says about *what is consistent with* training is silent about *what training actually selects*, and the gap between those is where most of the empirical field lives. Unprovenanced.

**Continuous and infinite spaces.** Every anchor is finite, and finiteness is load-bearing: exhaustive search, exact rationals, and counting arguments all depend on it. Chapter 2’s counting fact, Chapter 4’s grids and Chapter 6’s enumeration have no obvious analogue over continua without measure theory doing the work. Unprovenanced.

**Time, non-stationarity and feedback.** KW is one episode. Nothing here addresses a system deployed repeatedly into a world its own outputs change, and the wedge is at its worst exactly there – a proxy that was fine becomes bad *because* you optimised it. Chapter 3 measures optimisation pressure with a knob; in deployment the knob turns itself. Unprovenanced.

## The fence around the fences

An honest count: this tour verified 53 numbered claims on anchors that between them have no learning, no continuity, no time and one agent. What is on this side of the fence is exact. What that exactness is evidence *about* is the question every Honest scope section in the tour has tried to answer the same way, and the answer has never been “frontier models”.

# Part A – the anchor appendix

Every number printed anywhere in this tour has a receipt here, and every line below was generated by running the suite, not by transcription.

| # | Ch | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
|---|----|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1 | 1  | 4 policies, 3 behaviours, fibre sizes (2,1,1); the two off-policies coincide                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
| 2 | 1  | $V(\text{off})=1$ , $V(\text{on},a)=3/4$ , $V(\text{on},b)=0 \rightarrow$ off optimal, though r ranks A first                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| 3 | 1  | affine gauge exact at $\gamma=1$ over 400 draws; a bare constant shift flips the optimum at $\gamma=1/2$ (witness found by search)                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| 4 | 1  | 3 walls pairwise at 60 deg ( $\cos^2 = 1/4$ exactly); all 6 arcs found by search                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| 5 | 1  | 6 orderings $\rightarrow$ 3 behaviours, every fibre of size 2; top-ranked terminal always wins. At $\gamma=1$ that is exact for ANY distinct-valued reward, because $\text{value} = r[\text{outcome}]$ makes the top-ranked terminal optimal by definition – so the 200 draws per ordering from the integer pool -9..9 test the sampler, and the pool is not a scope condition. Nor is $\gamma=1$ the scope: swept exhaustively over all 5814 distinct-valued triples on that pool the ordering still determines the behaviour at $\gamma$ in $\{1, 99/100, 9/10\}$ , and stops doing so at $\{3/4, 1/2\}$ |
| 6 | 1  | same ordering ('B', 'bot', 'A'), two behaviours at $\gamma=1/2$ ; zero splits at $\gamma=1$ over the identical search                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| 7 | 2  | $P[\text{off optimal}] = 1/3$ by 1716-point solver sweep – every distinct-valued assignment of the 13-value integer pool -4..8 to the 3 terminals, $13 \times 12 \times 11$ ; rank-count agrees but is label-blind (all 3 terminals return $1/3$ ), so it ties exchangeability only                                                                                                                                                                                                                                                                                                                        |
| 8 | 2  | $1/(k+1)$ for $k=1..8$ by rank count; the solver route agrees for $k=1..5$ but on its default pool $\{0..k\}$ it sweeps only the $(k+1)!$ bijections from terminals to values (2 points at $k=1$ , 720 at $k=5$ ) – the same enumeration as the rank count under $\text{value} > \text{rank}$ , so here it re-labels rather than independently confirms; check 7's wider pool is where the two routes genuinely differ; $k=1$ gives exactly $1/2$ – bias vanishes                                                                                                                                          |
| 9 | 2  | $\text{swap}(\text{off})$ is a proper subset of on: $572 < 1144$ , 572 witnesses left over; $k=1$ control ties exactly                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |

| #  | Ch | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
|----|----|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 10 | 2  | depth-equalised: exactly $1/3$ at every gamma and every pool; real KW depths at gamma= $1/2$ give $9/14$ / $131/286$ / $17/126$ (pos/mixed/neg); the pools are the integer ranges 1..9, -4..8 and -9..-1 – so the middle one spans 13 values where the others span 9, and is mixed-sign rather than symmetric about 0 – enumerated over distinct-valued ordered triples and counting only a strict unique bot optimum                                                                                                                                                              |
| 11 | 2  | exchangeable prior gives exactly $1/3$ ; shifting bot's prior by +2 makes shutdown the majority outcome ( $5/9$ ); the prior is t0 and t1 uniform on the 10 CONSECUTIVE integers 0..9 with bot uniform on that grid shifted by +c, distinct-valued triples only, at gamma=1 on equalised depths; the +2 reported is the SMALLEST c in 0..11 exceeding $1/2$ , not the only one, and the step matters as much as the span – on a step-2 grid of 10 values the first majority shift is +4, also at $5/9$                                                                             |
| 12 | 3  | phi vanishing on terminals: optimal set unchanged on 512 of 600 attempted draws – the other 88 were dropped unexamined by the distinct-reward filter, so 'draws' is not the loop bound; both topologies, integer rewards on -9..9, integer non-terminal potentials on -6..6, zero on terminals, 4 discounts $\{1, 1/2, 1/3, 9/10\}$                                                                                                                                                                                                                                                |
| 13 | 3  | constant terminal potential: 146/3403 flips on KW, 0/3383 on depth-equalised KW -> the path-length channel is real and isolated; denominators are the draws that survived the distinct-reward filter out of 4000 attempted per topology (rewards three distinct integers on -9..9; terminal potential one non-zero constant on -6..6, so the interior potentials cancel; gamma uniform on $\{1/2, 1/3, 9/10\}$ – gamma=1 is excluded because the path-length term vanishes there and no undiscounted draw could flip, which is a different menu from check 14's)                   |
| 14 | 3  | varying terminal potential: 718/3415 flips on depth-equalised KW -> a second, independent failure channel; the denominator is the draws surviving the distinct-reward filter out of 4000 attempted (rewards distinct on -9..9; every terminal potential drawn independently on -6..6; gamma uniform on $\{1, 1/2, 9/10\}$ ). That menu is NOT check 13's $\{1/2, 1/3, 9/10\}$ , and gamma=1 is exactly where a constant terminal potential is inert while a varying one is not, so the two flip counts printed side by side are not a like-for-like comparison of the two channels |
| 15 | 3  | non-potential shaping term found by search that moves the optimum                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |



| #  | Ch | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|----|----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 16 | 3  | 6 proxy orderings: worst regret 11, reached while still agreeing on 1 of 3 comparisons; inversion count 1 costs $\{0,1\}$ and count 2 costs $\{1,11\}$ , so agreement rate does not control regret. True values are off=0, tea=1, flood=-10 on the depth-equalised anchor, so 11 is exactly the spread $\max(\text{true}) - \min(\text{true})$ and is fixed by that table rather than discovered; what transfers is the $\{1,11\}$ split at equal inversion count, not the magnitude                                                                                                                                                                                                                   |
| 17 | 3  | proxy rises at every step; truth peaks at $n=4$ (31.031) then falls to 4.313 at $n=25$ ; aligned control never turns over; over 200 random tables the curve turns over 200/200 times but the rise is visible in only 182 – when the decoupled item is bad enough the peak sits at $n=1$ . The table is 49 aligned items with truth = proxy = 0.48 plus one decoupled item at proxy 60 and truth -60, and each best-of- $n$ draws $n$ of those 50 i.i.d. uniformly WITH replacement; the two magnitudes are what put the peak at $n=4$                                                                                                                                                                  |
| 18 | 3  | 213/729 reward pairs unhackable on the full simplex – reward values drawn independently from $\{0,1,2\}$ at each of the 3 terminals, so $3^3 \times 3^3 = 729$ ordered (proxy, true) pairs with repeats allowed; policy search and the GRID-FREE centred-proportionality criterion agree exactly, which is what licenses reading a 25-point (p,q) search grid as the full simplex – the grid is the witness, not the domain. Non-constant survivors are precisely positive affine images: 60 of the 213 have both rewards non-constant, 36 of those with the two rewards distinct                                                                                                                      |
| 19 | 4  | cone closed under positive scaling, constants and addition; 572/572 negative scalings leave it; cross-section is exactly 2 of Chapter 1's 6 arcs. Members are the 572 of 1716 DISTINCT-VALUED triples over a 13-value integer grid (permutations, so this check contains no ties at all) that have A maximal – which makes A strictly MINIMAL in -r, so negation always exits and 572/572 is forced by construction rather than discovered. The rewards that do survive negation are the tied and constant ones this grid excludes, and they are exactly check 20's lineality space. The load-bearing parameter is the 13 values, not the endpoints: the counts do not move when the grid is recentred |

---

| #  | Ch | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|----|----|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 20 | 4  | lineality space is exactly the 11 constant vectors – one per value of the 11-value pool -5..5 – in this check’s grid of all $11^3$ tuples with repeats allowed; the distinct-valued grid of checks 19 and 21 contains no constant vector at all, which is why this check needs its own grid; the constant reward lies in all 3 cones simultaneously and makes every policy optimal                                                                                                                                                                                                                                                                                                                          |
| 21 | 4  | observing the argmax on all 4 sub-MDPs yields exactly 6 signatures = the 6 orderings; every signature class still contains gauge-inequivalent rewards                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| 22 | 4  | Luce demonstrator: 125 weight vectors -> 91 distinct choice distributions, every collision a pure rescaling; the noiseless observer has a vocabulary of exactly 7 observations in total, ties included; weights drawn independently from the pool $\{1,2,3,4,6\}$ at $\beta=1$ , $5^3 = 125$ vectors over the 3 terminals. The 91 is pool-dependent but $\beta$ -INdependent, and both are measured here rather than assumed: $\beta$ in $\{1,2,3\}$ all return 91, while the pool $\{1,2,3,4,5\}$ returns 115 and $\{1,2,4,8,16\}$ returns 61 – a pool with MORE internal ratios collides more proportional triples and so yields FEWER distinct distributions, which is the opposite of the natural guess |
| 23 | 4  | ( $\beta=2, r$ ) and ( $\beta=1, 2r$ ) give identical behaviour on all 64 weightings, the $4^3$ tuples over the pool $\{1,2,3,5\}$ ; fixing $\beta$ leaves only 3 collisions on that pool, and they are exactly the 4 constant vectors collapsing onto the uniform distribution – so ‘the fibre is rescalings alone’ is nearly empty here, and the non-trivial form of the claim is check 22’s $125 \rightarrow 91$                                                                                                                                                                                                                                                                                         |
| 24 | 4  | maximising $r$ and minimising $-r$ are identical on all 125 weightings, the $5^3$ tuples over the pool $\{1,2,3,5,7\}$ ; the agreement is an algebraic identity of the reciprocal parameterisation rather than a search – it holds for every non-zero rational weighting, so this is a demonstration and not evidence. The +1 perturbation control shows only that the comparison can discriminate at all, separating 120 of the 125 and failing exactly on the 5 constant tuples                                                                                                                                                                                                                           |

---

| #  | Ch | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|----|----|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 25 | 5  | wait $\geq \max(\text{act}, \text{off})$ on 600 beliefs; 575 strict and 25 tight, and tightness coincides exactly with $U$ being single-signed. Each belief is 5 distinct integer utilities drawn WITHOUT replacement from $-8..8$ , with weights uniform on $1..12$ and normalised – a spanning family, not a quantifier, and both counts are properties of it                                                                                                                                      |
| 26 | 5  | deference gap equals EVPI exactly on 400 beliefs, by two code paths; zero when the sign is known                                                                                                                                                                                                                                                                                                                                                                                                     |
| 27 | 5  | on 292 mixed-sign beliefs the deference threshold is $p^* = \max(E[U+], -E[U-])/E[\text{abs } U]$ , always at least $1/2$ , and a 240-point grid scan lands within one grid step of it                                                                                                                                                                                                                                                                                                               |
| 28 | 5  | over-confident belief loses the full EVPI on all 289 mixed-sign beliefs; over-dispersed belief loses exactly 0 – the anchor cannot reproduce the source’s second failure mode                                                                                                                                                                                                                                                                                                                        |
| 29 | 5  | over 40 interruption rates Q-learning stays indifferent and Sarsa strictly prefers the uninterruptible route; at $q=0$ both agree, so the test discriminates                                                                                                                                                                                                                                                                                                                                         |
| 30 | 6  | training-optimal set is a cone under positive scaling; its lineality is exactly the plane $w_1+w_2+w_3=0$ , so training pins 1 of 3 dimensions and leaves the rest free                                                                                                                                                                                                                                                                                                                              |
| 31 | 6  | 98 of 153 training-optimal objectives fail at test (55/153 generalise); adding the decorrelated configuration to training removes every failure                                                                                                                                                                                                                                                                                                                                                      |
| 32 | 6  | a train/test distinguisher multiplies the consistent set by exactly 343 – the size of the hypothesis box $(-3..3)^3$ , 7 integer weights per feature – and drives the generalising fraction to the prior 153/343 (153 of the 343 box points prefer the true feature at test), i.e. training carries zero information about test behaviour; without the flag it carries 55/153 (55 of the 153 strictly train-optimal points), which is WORSE than the prior – correlated training actively misinforms |
| 33 | 6  | over the 343-point integer box $(-3..3)^3$ the training-indistinguishable directions form a subspace: 37 grid points on one comparison – the plane $w_1+w_2+w_3=0$ , dimension 2 – and 7 on two – the line $w_1=0, w_2=-w_3$ , dimension 1. The two counts are lattice-point counts of that plane and that line and move with the box; the dimension drop from 2 to 1 does not. Same cone-and-lineality structure as Chapter 4, with the training set doing the observing                            |

| #  | Ch | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
|----|----|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 34 | 6  | consistent objectives fall 153 -> 55 -> 37 -> 19 over the 343-point weight box $(-3..3)^3$ , as the comparisons 111>000, 100>011, 110>001, 101>010 are imposed cumulatively in that order, each as a strict preference; every survivor of full coverage has positive weight on the true feature                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
| 35 | 7  | BT likelihood exactly invariant under 400 rescalings of the weights, and moved by a non-uniform bump; in score space this is the additive constant of Chapters 1 and 4                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| 36 | 7  | 8 one-per-pair datasets: 6 transitive and unidentifiable, 2 cyclic and identifiable – only the data that looks inconsistent pins a finite score; and (0, 1, 2) is an n=2 dataset the singleton-only test wrongly accepts                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
| 37 | 7  | separable data: likelihood strictly climbs over 24 doublings and stays below 1; cyclic data: likelihood falls off in all 6 boundary directions, so the maximum is interior                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |
| 38 | 7  | KL pressure knob k: proxy rises at every step, truth peaks at k=1 then falls monotonically; aligned control never turns over; the curve turns over in 25/25 random tables – Chapter 3’s wedge from a different lever. The table is 25 aligned outcomes with $u = j+1$ and $\text{truth} = j$ , plus one decoupled flood at $u = 28$ and $\text{truth} = 40$ , under a uniform reference over all 26, swept $k = 0..9$ ; the aligned control is a SEPARATE table (uniform on 25, $\text{truth} = u = j+1$ ), not this one with the flood removed. And the random tables certify only THAT the curve turns over: over the whole 8 x 61 family ( $u\_flood\ 26..33$ , $t\_flood\ -80..-20$ ) the peak sits at $k=0$ in 86 cases, $k=1$ in 366 cases, $k=2$ in 36 cases – so $k=1$ is this table’s peak, not the family’s |
| 39 | 7  | preference implied by the reward equals the one implied by the policy ratio on 300 draws, with Z cancelling identically; 50 of 200 sampled references had zero support, where the inversion is undefined – that 50 is a sampling artefact and depends on how much of the shared stream the earlier checks consumed, so it is not a finding. Each component is drawn uniformly from $\{0..9\}/20$ , and the exact count over that whole grid is 271 of 1000, i.e. $10^3 - 9^3$ ; the content is the witness that the inversion is undefined there, not either integer                                                                                                                                                                                                                                                  |
| 40 | 8  | leaves/debate-path ratio $2^d/d$ rises 2 -> 1024/3 over depths 1..12; it ties at $d=1$ and $d=2$ and is strictly increasing after                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |

| #  | Ch | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|----|----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 41 | 8  | over all 6561 depth-3 trees – every assignment of leaf values from $\{0,1,2\}$ to the 8 leaves, $3^8$ – and 3 overclaims each, the true max plus 1 / 2 / 5: honesty wins every time under a correct judge and loses every time under an inverted one                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| 42 | 8  | to keep overall reliability at 99% the per-query judge accuracy must rise from 99/100 at depth 1 to 1249/1250 at depth 12 – the depth discount is paid for at the judge                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| 43 | 8  | with every leaf checkable honesty wins all 256 depth-2 positions – the $4^4$ assignments of leaf values from $\{0,1,2,3\}$ to the 4 leaves, each against an overclaim of the true max plus 1, a pool one value wider than check 41's – and making any single leaf unverifiable forces a draw in all 1024 cases                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
| 44 | 8  | mean PGR falls from 0.958 (positive in 12/12 trials) at 1 wrong label of 32 to 0.691 (12/12) at 12 of 32 – the student beats its supervisor at low noise and sometimes recovers nothing at high noise; a search found a 10-error pattern recovering exactly nothing. Design: the dataset is 32 labelled comparisons; student hypotheses are the 343-point box $(-3..3)^3$ ; each mean is over 12 flip patterns drawn from the shared seed, so these are Monte-Carlo point estimates that depend on the whole suite running in order, not exact quantities – exhaustively over all 32 one-error patterns the mean is 39/40. w2s accuracy is averaged over the entire argmax tie set (19-20 of the 343 weights tie at one error), and under a worst-case tie-break 8 of those 32 patterns give PGR = -1, so the sign of the low-noise effect is a property of the averaging rule |
| 45 | 9  | frame vectors sum to zero, norms 2/3, pairwise cosine exactly -1/2 in rationals                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
| 46 | 9  | over all 2024 triples of 24 rational planar directions the largest pairwise $\cos^2$ is never below 1/4 and never equals it – the best is 9/34; the projected frame attains 1/4 exactly, which is why the construction is a projection                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| 47 | 9  | superposition beats dedicating dimensions at every sparsity below 91/200 and never at or above $23/50 = 92/200$ – ADJACENT points of the 199-point grid $i/200$ , so the pair brackets the crossover rather than equalling it; the crossover is the root of $3p^2+3p-2$ , pinned by an exact sign change, and refining to $i/1000$ narrows the bracket to 457/1000..229/500                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |

| #  | Ch    | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
|----|-------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 48 | 9     | exhaustive over all 17296 unordered triples of the 48 non-zero integer directions in $(-3..3)^2$ : 80 have a common pairwise $\cos^2$ other than $1/4$ , every one of them 1; 8 of the 80 DO sum to zero, so sum-to-zero alone does not force $1/4 - 0$ have equal norms, and equal norms is the condition that does                                                                                                                                                                                                                                                            |
| 49 | 9     | 8 activation patterns collapse to 7 codes; the single collision is exactly $\{\}$ against $\{1,2,3\}$ – silence and saturation are the same vector                                                                                                                                                                                                                                                                                                                                                                                                                              |
| 50 | 10    | 6 protocols: exactly 1 strictly dominated; the remaining 5 form a strict frontier from (useful 1, safe $1/2$ ) to (useful 0, safe 1); suspicion scores are 0..4 with benign mass ( $1/2$ , $1/4$ , $1/8$ , $3/32$ , $1/32$ ) and attack mass its reverse, and safety is the attack mass at or above $\min(t, 4)$ – deferral plus an always-audited top score. So the safe- $1/2$ endpoint is exactly the chosen parameter ATTACK[4], while (useful 0, safe 1) is forced by total deferral for any normalised pair; the audited benign rate is $1/32$ , inside the $1/20$ budget |
| 51 | 10    | a colluding monitor drives safety to exactly 0 at every threshold but total deferral, against $3/4$ for a trusted one at deferral threshold $t=3$ SPECIFICALLY; the trusted baseline is threshold-dependent under the same score distributions, running from $31/32$ at $t=1$ to $1/2$ at $t=4$ , so $3/4$ is one point on that range and not a property of trusted monitors; a single synthetic probe detects the collusion with certainty                                                                                                                                     |
| 52 | 10    | five independent layers cut failure from $1/2$ to $1/32$ ; at correlation 1 the same five layers cut it to $1/2$ , and the benefit of depth falls monotonically to exactly zero as correlation rises                                                                                                                                                                                                                                                                                                                                                                            |
| 53 | 10    | ordinary adversarial training thins the backdoored population $83 \rightarrow 37 \rightarrow 15$ – a cut of 82% that never reaches zero; every one of the 15 survivors still fires on the untouched trigger, and only training on the trigger removes them, which requires knowing it                                                                                                                                                                                                                                                                                           |
| 54 | audit | 10 numbered chapters, each with exactly 1 key claim, 1 signpost and 1 demystification; 10 starred sections with none; receipts referenced gapless 1..53; with a prose mention injected a naive phrase count says 11 and the tight opener still says 10; part letters ABCDEF contiguous from A                                                                                                                                                                                                                                                                                   |
| 55 | audit | 45 rows non-decreasing; 7 era bands contiguous and covering with 3 ambiguous seams; 41 life spans printed and 12 left blank; 5 written/published pairs, widest gap 53 years                                                                                                                                                                                                                                                                                                                                                                                                     |

| #  | Ch  | What the check establishes                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
|----|-----|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 56 | c2c | swept all 625 shaping terms: optimality-preserving IF AND ONLY IF realisable by a terminal-vanishing potential, 13 of 625, zero mismatches – the converse is now computed on the anchor, not cited; at $\gamma=1/2$ , with shaping values on the 5-value grid -2..2 on each of the 4 edges ( $5*5*5*5$ ) and rewards all 210 distinct-valued ordered triples from -3..3; the surviving count depends on $\gamma$ , so $\gamma$ is part of the claim and not a detail                                                                                                                                                                                                                                                                                                           |
| 57 | c2c | every one of 391 distinct behaviours in a 6-planner by 216-reward class admits at least 2 decompositions, and up to 36; no behaviour identifies its pair. The rationality confound is only partly reachable: a beta-partner needs $w^{\wedge\beta}$ , which leaves any finite pool. The planners are Luce with exponent $\beta$ in $\{+/-1, +/-2, +/-3\}$ , negative meaning the minimiser $u = w^{\wedge-\beta}$ ; the 216 reward weights are all $6^3$ tuples over the pool $\{1,2,3,4,6,12\}$ , chosen CLOSED under $x \rightarrow 12/x$ so that a minimiser's reciprocal partner rescales back into the grid – on the un-closed pool $\{1,2,3,4,5,6\}$ the same sweep leaves 432 singleton fibres and the 'at least 2' fails; rewards are $r = \log w$ . The maximum 36 is |
| 58 | c2c | Q-learning indifferent and Sarsa strictly avoidant across 180 (reward, rate) combinations, with the $q=0$ control tying in every one                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
| 59 | c2c | over 343 grid points, separable data has 0 local maxima and cyclic data has 7, identically in all three step regimes and under a gentler step – one direction suffices here and the agreement is measured, not assumed; a degenerate step of ratio 1 calls all 343 points maxima, so the search is not vacuous; the grid is the 7 powers of 2 from $1/8$ to 8 on each of 3 weights ( $7*7*7$ ), and 'local' means no strict improvement from multiplying one weight by $3/2$ or its reciprocal – a probe finer than the grid spacing, landing strictly between grid points, with ties not counting as improvements, which is what makes the ratio-1 control read 343                                                                                                           |

All 59 checks pass, alongside 4 standing audits. The suite uses exact rational arithmetic, exhaustive enumeration and searched pools throughout; no result depends on a floating-point tolerance anywhere, and the **converges** helper shipped for that purpose is declared unused rather than deleted. Checks 56 to 59 are the cited-to-computed pass and are marked c2c.

## Part B – what the field repeats, and what the sources say

Two research passes produced a set of corrections that would otherwise have lived only in comments the reader never sees. Verified but invisible is a delivery failure, so they are printed here.



| What is repeated                                          | What the source says                                                                                       | Verdict           |
|-----------------------------------------------------------|------------------------------------------------------------------------------------------------------------|-------------------|
| Goodhart wrote "when a measure becomes a target..."       | Goodhart 1975 says an observed regularity collapses under control pressure; the sentence is Strathern 1997 | MISATTRIBUTED     |
| Wiener's sentence is a paraphrase                         | it is verbatim in Science 131(3410), with the bottle factory, the broom, the genie and the monkey's paw    | CONFIRMED         |
| Wiener invokes Goethe                                     | he does not name Goethe anywhere in the article                                                            | FOLKLORE ADDITION |
| The 2015 paper's authors coined "corrigibility"           | the term was suggested by Robert Miles, who is not an author                                               | MISATTRIBUTED     |
| "Optimal policies tend to seek power" is Turner's         | it is a five-author paper: Turner, Smith, Shah, Critch and Tadepalli                                       | COMPOSITE         |
| Bradley and Terry invented the model in 1952              | Zermelo published the same structure in 1929                                                               | PRIOR ART         |
| Q-learning and Sarsa are both safely interruptible        | Q-learning already is; Sarsa is not without modification                                                   | GARBLED           |
| DPO is "Rafailov, Sharma, Mitchell, Ermon, Manning, Finn" | the published order puts Manning before Ermon                                                              | GARBLED           |
| Omohundro lists four or five AI drives                    | the paper prints no headline count; a six-item reading fits the text                                       | GARBLED           |

| What is repeated                                          | What the source says                                                                            | Verdict             |
|-----------------------------------------------------------|-------------------------------------------------------------------------------------------------|---------------------|
| Omohundro is at pp. 483-493                               | the ACM record gives pp. 483-492                                                                | CORRECTED           |
| Superposition does not connect to compressed sensing      | Toy Models has a related-work section on it and names Candes and Tao                            | REFUTED             |
| Toy Models states the 120-degree, minus-one-half geometry | it states the triangle and the Thomson connection; the numbers are consequences, not quotations | OVERREAD            |
| "AI safety via debate" is a published paper               | arXiv and a blog post; never peer-reviewed                                                      | STATUS              |
| "Risks from learned optimization" is a published paper    | arXiv only, no venue                                                                            | STATUS              |
| The Amodei on Concrete Problems is Daniela                | it is Dario on both that and deep RL from human preferences                                     | WRONG SIBLING       |
| Ridgway's given name is known                             | one database's expansion is algorithmic; the byline is initials only                            | NOT ESTABLISHED     |
| "Inner alignment" pre-dates 2019                          | two passes searched and neither found a citable earlier use                                     | PERSISTENT NEGATIVE |

## Counts that are one step from each other

Five concrete problems. Five convergent instrumental values. Four Goodhart variants. Six Omohundro drives on the reading this tour uses. Nine gridworlds, on secondary evidence only. These are small numbers about adjacent topics and they are easy to swap; each was settled by asking rather than by recall.

## How often a fact was looked at

A flag records what kind of evidence a fact has. It does not record how much. Four citations in this tour – Goodhart 1975, Strathern 1997, the reward-hacking paper's pagination and the power-seeking paper's pagination – were confirmed by one pass and not re-reached by later ones. That is single-pass verification, and it is not the same as double. The second pass recommended demoting them to memory-flagged on those grounds; this tour declined, because absence of evidence is not evidence of absence, and a flag that means "verified most recently" rather than "verified" would be worse than useless.

# Part C – a timeline, with its blanks

This part had its own research pass. Contribution dates are not life dates, and a timeline assembled from a bibliography is a bibliography. The rows below are generated from the same table check 55 verifies structurally, so a transposed row cannot ship silently.

## Era bands

Bands are interpretive. Three seams are genuinely contested and are marked as such rather than hidden behind a clean line.

| Band                                | Years     | Seam at its start |
|-------------------------------------|-----------|-------------------|
| Foundations and cybernetics         | 1927-1964 | clean             |
| Dormancy                            | 1965-1999 | AMBIGUOUS         |
| Pre-institutional modern era        | 2000-2013 | clean             |
| Field consolidation                 | 2014-2015 | AMBIGUOUS         |
| Deep-learning safety                | 2016-2021 | clean             |
| Deployment inflection               | 2022-2022 | AMBIGUOUS         |
| Governance, evaluations and control | 2023-2026 | clean             |

The two that matter most: whether the field begins in the pre-institutional 2000s or with the 2014-2016 consolidation, and whether late 2022 is a genuine boundary or merely a change in public attention. This tour takes a position on neither.

## The rows

Where a work was written or delivered materially before it was published, both dates are printed. One date would lie by omission.

| Published | Written | Entry                                                                      | Flag |
|-----------|---------|----------------------------------------------------------------------------|------|
| 1927      |         | Thurstone, law of comparative judgement                                    | S    |
| 1929      |         | Zermelo, ranking as a maximum-likelihood problem                           | P    |
| 1938      |         | Samuelson, revealed preference (with an August addendum)                   | P    |
| 1942      |         | Asimov's laws, as a cultural marker                                        | M    |
| 1947      | 1944    | von Neumann and Morgenstern, axiomatic utility<br>ADDED in the 2nd edition | S    |
| 1952      |         | Bradley and Terry, paired comparisons                                      | P    |
| 1954      |         | Olds and Milner, self-stimulation in rats                                  | S    |
| 1956      |         | Ridgway, dysfunctional consequences of performance measurement             | P    |
| 1957      |         | Ford Jr, existence condition for the paired-comparison MLE                 | P    |

| Published | Written | Entry                                                                            | Flag |
|-----------|---------|----------------------------------------------------------------------------------|------|
| 1960      |         | Wiener, Science 131(3410); Samuel's refutation, Science 132(3429)                | P    |
| 1965      |         | Good, speculations on the first ultraintelligent machine                         | P    |
| 1967      |         | Afriat, construction of utility functions from expenditure data                  | P    |
| 1969      | 1948    | Turing, Intelligent Machinery (NPL report; publication year disputed)            | S    |
| 1972      |         | James P. Anderson, Computer Security Technology Planning Study                   | S    |
| 1974      |         | Welch, lower bounds on maximum cross-correlation                                 | P    |
| 1975      |         | Goodhart, Problems of Monetary Management; Kerr, on the folly of rewarding A     | P    |
| 1975      |         | Saltzer and Schroeder, eight design principles – defence in depth NOT among them | P    |
| 1985      |         | Goldwasser, Micali and Rackoff, STOC; Babai, STOC                                | P    |

| Published | Written | Entry                                                                                                           | Flag |
|-----------|---------|-----------------------------------------------------------------------------------------------------------------|------|
| 1989      | 1985    | Goldwasser, Micali and Rackoff, journal version                                                                 | P    |
| 1990      |         | NIH designates thirteen species as model organisms                                                              | S    |
| 1992      | 1990    | Shamir, $IP = PSPACE$ ; Lund, Fortnow, Karloff and Nisan                                                        | P    |
| 1997      |         | Strathern, the sentence everyone attributes to Goodhart                                                         | P    |
| 1998      |         | Russell, COLT: the inverse problem posed; Network Security Framework 1.0 brings defence in depth into computing | P    |
| 2000      |         | Ng and Russell, algorithms for inverse RL; SIAI founded                                                         | P    |
| 2003      | 1951    | Bostrom, ethical issues in advanced AI (the paperclip example)                                                  | S    |
| 2004      |         | Turing, Intelligent Machinery: A Heretical Theory, first widely published                                       | S    |
| 2005      |         | Future of Humanity Institute founded                                                                            | S    |

| Published | Written | Entry                                                                                                 | Flag |
|-----------|---------|-------------------------------------------------------------------------------------------------------|------|
| 2006      |         | Donoho; Candes, S<br>Romberg and Tao:<br>compressed sensing                                           |      |
| 2008      |         | Omohundro, The Basic P<br>AI Drives, pp. 483-492;<br>Ziebart et al., MaxEnt<br>IRL                    |      |
| 2009      |         | Erhan et al., visualis- S<br>ing higher-layer features;<br>LessWrong launched                         |      |
| 2012      |         | Bostrom, The Superintel- S<br>ligent Will                                                             |      |
| 2013      |         | SIAI renamed the Ma- P<br>chine Intelligence Re-<br>search Institute                                  |      |
| 2014      |         | Bostrom, Superintelli- S<br>gence; Future of Life<br>Institute founded                                |      |
| 2015      |         | Puerto Rico conference (2- P<br>5 Jan); Corrigibility; Ope-<br>nAI announced (11 Dec)                 |      |
| 2016      |         | Concrete Problems; P<br>CIRL; Safely Inter-<br>ruptible Agents; CHAI<br>founded                       |      |
| 2017      |         | Off-Switch Game; deep P<br>RL from human prefer-<br>ences; gridworlds; Asilo-<br>mar; Distill founded |      |

| Published | Written | Entry                                                                                                                      | Flag |
|-----------|---------|----------------------------------------------------------------------------------------------------------------------------|------|
| 2018      |         | Armstrong and Minder-<br>mann; AI safety via de-<br>bate; Goodhart taxon-<br>omy; Alignment Forum                          | P    |
| 2019      |         | Hubinger et al., risks<br>from learned optimization<br>(arXiv only)                                                        | P    |
| 2020      |         | Zoom In: circuits; obfus-<br>cated arguments; learning<br>to summarize                                                     | P    |
| 2021      |         | Anthropic incorporated;<br>ARC founded; Distill hia-<br>tus; Transformer Circuits<br>begun                                 | P    |
| 2022      |         | Toy Models of Super-<br>position (14 Sep); re-<br>ward hacking; Instruct-<br>GPT; red-teaming LMs                          | P    |
| 2023      |         | CAIS statement (30 May);<br>Bletchley (1-2 Nov); EO<br>14110 (30 Oct); Towards<br>Monosemanticity                          | P    |
| 2024      |         | AI Control (ICML);<br>Sleepers Agents; weak-to-<br>strong; Scaling Monose-<br>manticity; FHI closed;<br>EU AI Act in force | P    |
| 2025      |         | UK institute renamed to<br>AI Security (14 Feb);<br>Paris summit; SAE cri-<br>tiques accumulate                            | S    |
| 2026      |         | Sparse-autoencoder sanity<br>checks report low true-<br>feature recovery                                                   | S    |

## Life dates, and the blanks

41 spans are printed below. 12 names are printed blank, and the blank is the honest output: three research passes found no birth year for them in any reliable public source. A guess next to a name is a fabrication, and this tour would rather print nothing. The asymmetry is itself a fact about the field – its historical figures are documented and its living researchers are not.



| Name        | Born | Died |
|-------------|------|------|
| Wiener      | 1894 | 1964 |
| Samuel      | 1901 | 1990 |
| von Neumann | 1903 | 1957 |
| Morgenstern | 1902 | 1977 |
| Savage      | 1917 | 1971 |
| Campbell    | 1916 | 1996 |
| Bellman     | 1920 | 1984 |
| Goodhart    | 1936 |      |
| Strathern   | 1941 |      |

| Name       | Born | Died |
|------------|------|------|
| Samuelson  | 1915 | 2009 |
| Houthakker | 1924 | 2008 |
| Varian     | 1947 |      |
| Thurstone  | 1887 | 1955 |
| Bradley    | 1923 | 2001 |
| Zermelo    | 1871 | 1953 |
| Luce       | 1925 | 2012 |
| Plackett   | 1920 | 2009 |
| Olds       | 1922 | 1976 |

| Name    | Born | Died |
|---------|------|------|
| Milner  | 1919 | 2018 |
| Thomson | 1856 | 1940 |
| Turing  | 1912 | 1954 |
| Good    | 1916 | 2009 |
| Jacobs  | 1863 | 1943 |
| Micali  | 1954 |      |
| Rackoff | 1948 |      |
| Babai   | 1950 |      |
| Shamir  | 1952 |      |

| Name       | Born | Died |
|------------|------|------|
| Fortnow    | 1963 |      |
| Sipser     | 1954 |      |
| Kozen      | 1951 |      |
| Stockmeyer | 1948 | 2004 |
| Russell    | 1962 |      |
| Ng         | 1976 |      |
| Bostrom    | 1973 |      |
| Omohundro  | 1959 |      |
| Yudkowsky  | 1979 |      |

| Name               | Born | Died |
|--------------------|------|------|
| Abbeel             | 1977 |      |
| Bengio             | 1964 |      |
| Barto              | 1948 |      |
| Anderson, Ross J.  | 1956 | 2024 |
| Anderson, James P. | 1930 | 2007 |

Printed blank: Ridgway, Kerr, Afriat, Terry, Christiano, Amodei, Olah, Krakovna, Leike, Hubinger, Shlegeris, Barnes.

### Name-collision boxes

A timeline puts surnames next to dates, which is exactly the context in which two people with one surname become one wrong person. Every collision below was checked and every one is real.

| Collision                                   | Resolution                                                                                                                                                                                        |
|---------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Ross J. Anderson / James P. Anderson</b> | two Andersons in computer security: Ross 1956-2024 wrote Security Engineering, James P. 1930-2007 wrote the 1972 Anderson Report. A timeline printing Anderson beside both dates would merge them |
| Stuart Russell / Stuart Armstrong           | two Stuarts, different work, frequently merged                                                                                                                                                    |
| Dario / Daniela Amodei                      | siblings; only Dario is an author on the papers cited here                                                                                                                                        |
| Six Andrews                                 | Ng, Maas, Bagnell, Critch, Lefrancq, Barto – all distinct                                                                                                                                         |
| Brown / Brown-Cohen / Everitt               | Tom B. Brown, Jonah Brown-Cohen, Tom Everitt – distinct                                                                                                                                           |
| Olah / van Merwijk / Manning                | three different Chrises                                                                                                                                                                           |
| Krueger / Manheim / Donoho                  | three different Davids                                                                                                                                                                            |
| Goodhart / Rackoff                          | two different Charleses                                                                                                                                                                           |
| Bostrom / Cammarata                         | two different Nicks                                                                                                                                                                               |

| Collision                      | Resolution                                                               |
|--------------------------------|--------------------------------------------------------------------------|
| L. R. Ford Junior / Senior     | the paired-comparison result is the son's                                |
| David Donoho = David L. Donoho | ONE person under two name forms; do not double-list                      |
| V. F. Ridgway                  | given name NOT established; the expansion in one database is algorithmic |

## What the tour records but does not build on

Named here so the timeline is a history rather than a reading list: Ford's existence condition and Welch's bound, both now cited in the chapters that use them; Turing's 1948 report and 1951 lecture; Good's 1965 intelligence-explosion paper; Asimov as a cultural marker; the founding of MIRI, FHI, FLI, CHAI, OpenAI, Anthropic, ARC and the safety institutes; the Puerto Rico, Asilomar, Bletchley, Seoul and Paris sequence; the venues – Distill, its hiatus, and the Transformer Circuits Thread that succeeded it; and the two imports this tour traces to their sources in Chapter 10.

Deliberately absent, because no pass could date them: Solomonoff induction, AIXI, the universal intelligence measure, approval-directed agents, the eliciting-latent-knowledge report, and Constitutional AI. They belong in a complete timeline and this one does not have them.

# Part D – the open frontier

What is proved, what is open, and what is merely believed. Every line carries the status its sources actually support, which in several cases is weaker than the status the field’s shorthand implies.

## Proved, with their hypotheses printed

**Power-seeking.** Certain graphical symmetries in a finite MDP make options-preserving behaviour optimal for most reward functions under a specified prior. The hypotheses are load-bearing and the authors press them harder than their readers do: the results concern **optimal** policies in **finite** MDPs under a **chosen distribution** over rewards, and optimal policies are, in the paper’s own framing, often qualitatively divorced from learned ones.

**Reward hacking.** Over the set of all stochastic policies, two reward functions can be unhackable only if one of them is constant or the two are positive affine transforms of one another – check 18 finds 60 pairs where both are non-constant, 36 of them with the two rewards distinct. Non-trivial unhackable pairs exist only over deterministic or finite policy sets. The mechanism is linearity of return in occupancy measures, and it is why Chapter 3’s result on the anchor is not a curiosity.

**Preference identifiability.** A policy admits no unique decomposition into a planner and a reward, and a simplicity prior does not select the true one. This is the tour’s synthesis theorem and it is the strongest negative result in the subject.

**Shaping.** Potential-based shaping preserves optimal policies, and the form is necessary as well as sufficient. The terminal-state condition is standard but neither research pass read it from the original, so it is carried as verified-secondary.

**Debate, conditionally.** With a polynomial-time errorless judge and computationally unbounded provers, optimal play decides PSPACE; cross-examination lifts it to NEXP. Both conditions are load-bearing, and the doubly-efficient line exists to remove the second by restricting provers to polynomial time.

## Open

**Obfuscated arguments.** A dishonest debater can construct an argument whose flaw neither side can locate within budget, forcing a draw. Unresolved. A 2025 prover-estimator variant attempts to make it a losing strategy, but only under an added stability assumption.

**Scalable oversight in general.** Weak-to-strong generalisation is an empirical result with limits its own authors name, including disanalogies they call important. There is no theorem.

**Interpretability completeness.** No result establishes that a circuit-level explanation is complete. Since 2024 a run of negative results has accumulated around sparse autoencoders – probes underperforming ordinary baselines, simple methods matching them at steering, interpretable-looking features recovered from randomly initialised transformers, and a 2026 sanity check reporting low true-feature recovery. One major laboratory publicly deprioritised the approach in 2025; a substantive counter-position holds that the tools are for discovering unknown concepts rather than acting on known ones.

**Control under capability growth.** The control results are empirical, in a specific coding testbed, with conservative red-team affordances, and the authors state that the guarantees degrade as the gap between untrusted and trusted models widens.

## The strongest honest statement available

**No result establishes that any named alignment technique scales to frontier models.** Every such claim in the literature is empirical and explicitly bounded, and the bounds are usually stated by the authors rather than by their critics. That sentence is the frontier.

## What this tour leaves open about itself

Four citations carry single-pass verification. A short list of facts remains memory-flagged after two passes. A third pass cleared four of them outright – the Erhan technical report number, Luce’s publisher and choice axiom, the DPO derivation’s support assumption, and a citation for Bradley-Terry shift-invariance. A fourth pass, given the three survivors and nothing else to do, closed all three: the page in Anderson, the debut of defence in depth in computing, and the coinage of model organism, the last of which is settled as a well-evidenced negative – there is no coiner, and the historians who looked say so. **So that list is now empty** – and a census run to confirm it found the sentence overclaiming, which is worth reporting rather than quietly fixing. **Exactly one flag survives that was never on that list:** the earliest clear statement of the evolution analogy, in Chapter 6, which no pass has dated. Two further items are recorded not as flags but as persistent negatives, which is a different thing: no pass found a pre-2019

use of inner alignment, and no coiner exists for model organism in biology. So the honest statement is not that the tour is free of error. It is that every fact in it has been looked at, the one that could not be settled is named, and the ones that came back empty are labelled empty rather than dressed as answers. Three fences are unprovenanced. And the largest gap is not a citation at all: the anchors have no learning, and everything the tour proves about what is consistent with training is silent about what training selects.

# Part E – the verification record

## Tallies, derived from the artifacts

| Quantity                                       | Value    |
|------------------------------------------------|----------|
| numbered checks, all passing                   | 59       |
| standing audits, all passing                   | 4        |
| data mutants killed                            | 40 of 40 |
| killed by named assertion                      | 40       |
| killed by crash (weaker evidence, M3)          | 0        |
| killed by timeout                              | 0        |
| survivors                                      | 0        |
| idempotent patches, three branches proven each | 35       |
| gates, all passing                             | 22       |

| Quantity                        | Value |
|---------------------------------|-------|
| pages, footered, envelope-clean | 85    |

## The cited-to-computed pass

Four claims the tour had marked cited, argued or instantiated were carried to a dedicated pass and asked one question: is this computable on the anchor?

| Claim                        |  | Outcome   | What changed                                                                                  |
|------------------------------|--|-----------|-----------------------------------------------------------------------------------------------|
| Shaping's converse           |  | CONVERTED | one searched instance became a biconditional over 625 terms, zero mismatches                  |
| Preference unidentifiability |  | PARTIAL   | two exhibited confounds became an exhaustive sweep of 391 behaviours, minimum fibre 2         |
| Interruptibility             |  | WIDENED   | one instance became 180 reward-and-rate combinations with the zero-rate control tying in each |
| Bradley-Terry non-existence  |  | CONVERTED | one path of doublings became a 343-point sweep finding no local maximum at all                |
| Debate's complexity results  |  | DECLINED  | quantifies over a complexity class; no finite anchor reaches it                               |

One full conversion, three partial, one deliberately declined. In every partial case the residual gap is named rather than implied: a quantifier over all MDPs, a rationality partner that leaves any finite pool, and a grid standing in for an orthant. **Knowing which is which was the point; converting everything was not.**

Gate mutations are not counted anywhere above, because weakening an assertion passes regardless and proves nothing. `mutate.py` refuses to run if any mutant edits an assertion, and that detector fired on the author twice.

## Defects, in three kinds

The three are separated because they have different causes and conflating them hides the pattern.

**Wrong constants – a claim's number was wrong.** The first KL-pressure sweep peaked at zero and produced five-hundred-digit rationals; rebuilt on small integers. The judge-reliability grid was too coarse to resolve the required accuracy below 1 at depth twelve. The Welch bound was asserted to be attained in a rational planar grid; it is not attained at all there, and the diagnosis became the chapter's best receipt.

**Wrong gates – a check could not have failed.** The construct-then-verify detector fired on its own whitelist, then on a legitimate initialiser, and was narrowed with a positive control rather than loosened. A Ford-condition clause was inert on the data it ran against. A proportionality helper's second cross-product was never exercised. An assertion quantified



over an empty slice and the tautology scan did not catch it; another was an always-true disjunction; a third passed on the wrong half of an **or**. The count audit matched a prose mention of a phrase rather than a box. A precondition – that the backdoor trigger lies outside adversarial training – was assumed rather than asserted. Two patch markers spanned line breaks and could never have been found after patching. **Every one of these was found by a mutant or by a gate built after an earlier instance of the same shape.**

**Wrong prose – the text claimed more than the computation showed.** Ten claims died in construction: that every damaging proxy inverts at least two comparisons; that an interior peak occurred in every random table; that the noiseless observer sees at most four outcomes; that the deference tightness condition is single-signedness, when zero is neutral; that a train/test flag lowers the generalising fraction, when it raises it and the real result is exact independence; that a leaves-to-queries ratio is strictly monotone, when it ties; that weak-to-strong recovers ground under any model of the supervisor; that the 2/3 result was confirmed by two independent routes, when one of them is label-blind; that adversarial training removes nothing, when it removes eighty-two per cent; and one report that a file was on disk when it was not.

## False negatives

The document has now been rendered, so the probe tally below is real rather than pending.

## The render phase, and its tally

**One false negative.** A booktabs-rule detector used a threshold of 500 pixels on a text block that is 125mm wide, which is 492 pixels at 100 dpi. It reported **zero tables in a document containing forty pages of them**. The threshold is now derived from the text width rather than typed, and validated against pages whose extracted text carries a known table header.

**Three false positives.** A reference probe matched the word "undefined" occurring in this document's own prose, which the log echoes back. A probe for a superseded page number fired on Part B's correction table, where the superseded number legitimately appears **as the thing being corrected**. A margin-note probe matched signpost wording reused in body prose and reported thirteen notes where there are ten. Each was narrowed and given a control; none was loosened.

**And one failure of a different kind, which belongs in this record because it is the worst sort.** Asked to inspect a rendered page, the author described a chapter opening. The extraction from that same page shows it is mid-chapter and carries two boxes. That was expected content narrated as seen. The renders at this resolution are not reliably legible,

so the visual channel was declared uninspectable and replaced throughout by objective detectors: box fill by a signed difference from paper against a sixty-seven-page no-box control, margin notes by ink in the note column against a control class with exactly zero, rules by a width-derived threshold, and structure by parsing the PDF independently of both the log and the source.

**The final proofread added two more probe failures, both mine.**

A contents check normalised the entry titles but not the pages it compared them against, and reported ten mismatches where there are none – apostrophes, commas and dashes. A column detector looked for continuous vertical whitespace and concluded that a five-column table had one column, because wrapped cells leave no continuous gutter. Both were replaced: the first by normalising both sides identically and running a deliberately wrong pairing as a control, the second by reading the real column specification out of the generated LaTeX and confirming the header cells appear in the render in order.

Running total after the final proofread: **two false negatives, four false positives, one false observation**. Not one of the seven was a defect in the document. Every one was a defect in the instrument pointed at it – and every visual claim in this document is bound to a phrase extracted from the page it describes.

## The final external round

A last verification pass was run against the finished PDF rather than against the source, because what shipped is what matters. It audited roughly fifty bibliographic assertions, fourteen folklore verdicts, forty-one printed life spans and seven event dates.

It raised **two alerts, and only one was live**. The live one was a wrong issue number for Bostrom 2012, printed as 22(1) where the correct citation is 22(2), pp. 71-85; it is corrected in the bibliography, and the superseded number now appears only inside its own correction. The second alert suspected that this document attributed the lift of debate to NEXP to the 2018 debate paper rather than to Barnes and Christiano. **It did not**. A search found zero sentences bundling the two, and the NEXP result was already carried as verified-secondary, separately from the primary-verified attribution. Patching would have damaged a correct passage.

That is two alerts triaged, one live and one false, against five false positives in ten alerts on a previous run. The discipline that matters is not the rate; it is that no alert is acted on before the exact printed string has been read.

The round also confirmed the two citations most likely to be flagged wrongly by a re-checker. The direct-preference-optimisation author order printed here matches the official proceedings, while one major index inverts it. And the warning that Candes, Romberg and Tao at 52(2) is a different paper from Candes and Tao at 52(12) is itself correct. Both would have been corrected into error by a pass that trusted secondary indices. What has occurred instead is five **false positives** from source-side scans: two

from the construct-then-verify detector firing on legitimate code, two from the tautology scan matching its own pattern table, and one from the count audit matching prose. Each was resolved by narrowing the pattern **and** binding a positive control to it, never by loosening.

## What the record cannot tell you

The suite and the document share machinery. A defect in that shared machinery corrupts both sides of every cross-tie identically and no check here would notice. This is the same limitation Chapter 10 records about correlated defences, and the tour is not exempt from its own result.

# Part F – primary bibliography and a reading path

Flags: P primary-verified, S verified-secondary, M not verified. Where two research passes disagreed or one could not re-reach a source, the entry says so.

## By chapter

**Chapter 1.** Wiener, Some Moral and Technical Consequences of Automation, *Science* 131(3410):1355-1358, 6 May 1960 [P]. Samuel, A Refutation, *Science* 132(3429):741-742, 1960 [P]. Ridgway, Dysfunctional Consequences of Performance Measurements, *Administrative Science Quarterly* 1(2):240-247, 1956 [P]. Kerr, On the Folly of Rewarding A While Hoping for B, *Academy of Management Journal* 18(4):769-783, 1975 [P]. von Neumann and Morgenstern, *Theory of Games and Economic Behavior*; axiomatic utility in the second edition of 1947 [S].

**Chapter 2.** Omohundro, The Basic AI Drives, AGI 2008, *Frontiers in AI and Applications* 171:483-492 [P]. Bostrom, The Superintelligent Will, *Minds and Machines* 22(2), pp. 71-85, 2012 [primary-verified; corrected by the final external pass from the 22(1) this document previously printed]. Turner, Smith, Shah, Critch and Tadepalli, Optimal Policies Tend to Seek Power, *NeurIPS* 2021, pp. 23063-23074 [P, single pass]. Turner and Tadepalli, Parametrically Retargetable Decision-Makers, *NeurIPS* 2022 [P].

**Chapter 3.** Goodhart, Problems of Monetary Management, *Papers in Monetary Economics* 1, Reserve Bank of Australia, 1975 [P, single pass]. Strathern, Improving Ratings, *European Review* 5(3):305-321, 1997 [P, single pass]. Ng, Harada and Russell, Policy Invariance Under Reward Transformations, *ICML* 1999, pp. 278-287 [P]; the terminal-state condition [S]. Manheim and Garrabrant, Categorizing Variants of Goodhart’s Law, arXiv 1803.04585 [P]. Skalse, Howe, Krashenninnikov and Krueger, Defining and Characterizing Reward Hacking, *NeurIPS* 2022, pp. 9460-9471 [P, single pass]. Olds and Milner 1954; Ring and Orseau, AGI 2011 [S].

**Chapter 4.** Samuelson, *Economica* 5(17):61-71 with the addendum at 5(19):353, 1938 [P]. Afriat, *International Economic Review* 8(1):67-77, 1967 [P]. Russell, Learning Agents for Uncertain Environments, *COLT* 1998,

pp. 101-103 [P]. Ng and Russell, Algorithms for Inverse Reinforcement Learning, ICML 2000, pp. 663-670 [P]; the degeneracy and cone results [S]. Ziebart, Maas, Bagnell and Dey, AAAI 2008, pp. 1433-1438 [P]. Armstrong and Mindermann, NeurIPS 2018, pp. 5603-5614 [P].

**Chapter 5.** Hadfield-Menell, Dragan, Abbeel and Russell, Cooperative Inverse Reinforcement Learning, NIPS 2016, pp. 3909-3917 [P]; The Off-Switch Game, IJCAI 2017, pp. 220-227 [P]. Orseau and Armstrong, Safely Interruptible Agents, UAI 2016, pp. 557-566 [P]. Soares, Fallenstein, Yudkowsky and Armstrong, Corrigibility, AAAI Workshop 2015 [P]; the term named by Robert Miles [S]. Leike, Martic, Krakovna, Ortega, Everitt, Lefrancq, Orseau and Legg, AI Safety Gridworlds, arXiv 1711.09883 [P].

**Chapter 6.** Hubinger, van Merwijk, Mikulik, Skalse and Garrabrant, Risks from Learned Optimization, arXiv 1906.01820, 2019, no venue [P]. Langosco, Koch, Sharkey, Pfau and Krueger, ICML 2022 [S]; Shah et al., arXiv 2210.01790 [S] – two distinct papers routinely cited as one.

**Chapter 7.** Thurstone, Psychological Review 34, 1927 [S]. Zermelo, Mathematische Zeitschrift 29(1):436-460, 1929 [P]. Bradley and Terry, Biometrika 39(3/4):324-345, 1952 [P]. Ford Jr, American Mathematical Monthly 64(8, Part 2):28-33, 1957 [P]. Christiano, Leike, Brown, Martic, Legg and Amodei, NIPS 2017, pp. 4299-4307 [P]. Ouyang et al., NeurIPS 2022, pp. 27730-27744, twenty authors [P]. Rafailov, Sharma, Mitchell, Manning, Ermon and Finn, Direct Preference Optimization, NeurIPS 2023 [S for order].

**Chapter 8.** Goldwasser, Micali and Rackoff, STOC 1985, pp. 291-304; SIAM Journal on Computing 18(1):186-208, 1989 [P]. Babai, STOC 1985, pp. 421-429 [P]. Shamir, JACM 39(4):869-877, 1992 [P]. Lund, Fortnow, Karloff and Nisan, JACM 39(4):859-868, 1992 [P] – adjacent papers, recorded together to prevent transposition. Irving, Christiano and Amodei, arXiv 1805.00899, 2018, no venue [P]. Christiano, Shlegeris and Amodei, arXiv 1810.08575, preprint [P]. Barnes and Christiano, obfuscated arguments, Alignment Forum 2020 [S]. Burns et al., ICML 2024 [P]. Brown-Cohen, Irving and Piliouras, arXiv 2311.14125 [P].

**Chapter 9.** Olah, Mordvintsev and Schubert, Distill 2(11), 2017, DOI 10.23915/distill.00007 [P]. Olah, Cammarata, Schubert, Goh, Petrov and Carter, Distill 5(3), 2020, DOI 10.23915/distill.00024.001 [P]. Elhage et al., Toy Models of Superposition, Transformer Circuits Thread, 14 September 2022, sixteen authors [P]. Welch, IEEE Transactions on Information Theory IT-20(3):397-399, 1974 [P]. Donoho, IEEE TIT 52(4):1289-1306, 2006 [S]; Candes, Romberg and Tao, IEEE TIT 52(2):489-509, 2006 [S] – not to be confused with Candes and Tao at 52(12):5406-5425. Bricken et al. 2023; Templeton et al. 2024 [S].

**Chapter 10.** Greenblatt, Shlegeris, Sachan and Roger, AI Control, ICML 2024, PMLR 235:16295-16336 [P]. Anderson, Security Engineering, three editions 2001, 2008, 2020 [S]. Shevlane et al., arXiv 2305.15324, twenty-one authors [P]. Phuong et al., arXiv 2403.13793, twenty-seven authors [P]. Hubinger et al., Sleeper Agents, arXiv 2401.05566, thirty-nine authors [P]. Perez et al., EMNLP 2022, pp. 3419-3448 [P].

## A reading path

If you read four things, read Wiener 1960 for the question, Ng, Harada and Russell 1999 for what a clean theorem in this subject looks like, Armstrong and Mindermann 2018 for the strongest negative result, and Greenblatt et al. 2024 for what the field does when the negative results are taken seriously.

If you read four more, add Omohundro 2008 for the argument before it was formalised, Hubinger et al. 2019 for the problem that has no theorem yet, Elhage et al. 2022 for the one place in the subject where a prediction about internal structure came out exactly right, and the obfuscated-arguments post for an open problem stated plainly by the people whose proposal it obstructs.

Read Strathern 1997 to see how a sentence acquires the wrong name.